

# Red-Black Trees


## Red-Black Trees

- *Red-black trees*:
  - Binary search trees augmented with node color
  - Operations designed to guarantee that the height  $h = O(\lg n)$
- First: describe the properties of red-black trees
- Then: prove that these guarantee  $h = O(\lg n)$
- Finally: describe operations on red-black trees

# Red-Black Properties

- The *red-black properties*:
  1. Every node is either red or black
  2. Every leaf (NULL pointer) is black
    - Note: this means every “real” node has 2 children
  3. If a node is red, both children are black
    - Note: can’t have 2 consecutive reds on a path
  4. Every path from node to descendant leaf contains the same number of black nodes
  5. The root is always black

Copyright © The McGraw-Hill Companies, Inc. Permission required for reproduction or display.



**Figure 13.1** A red-black tree with black nodes darkened and red nodes shaded. Every node in a red-black tree is either red or black, the children of a red node are both black, and every simple path from a node to a descendant leaf contains the same number of black nodes. (a) Every leaf, shown as a NIL, is black. Each non-NIL node is marked with its black-height; NIL's have black-height 0. (b) The same red-black tree but with each NIL replaced by the single sentinel  $nil[T]$ , which is always black, and with black-heights omitted. The root's parent is also the sentinel. (c) The same red-black tree but with leaves and the root's parent omitted entirely. We shall use this drawing style in the remainder of this chapter.

## Red-Black Trees

- Put example on board and verify properties:
  1. Every node is either red or black
  2. Every leaf (NULL pointer) is black
  3. If a node is red, both children are black
  4. Every path from node to descendent leaf contains the same number of black nodes
  5. The root is always black
- *black-height*: # black nodes on path to leaf
  - Label example with  $h$  and bh values

## Height of Red-Black Trees


- *What is the minimum black-height of a node with height  $h$ ?*
- A: a height- $h$  node has black-height  $\geq h/2$
- Theorem: A red-black tree with  $n$  internal nodes has height  $h \leq 2 \lg(n + 1)$

## RB Trees: Worst-Case Time

- We will prove that a red-black tree has  $h = O(\lg n)$  height
- Corollary: These operations take  $O(\lg n)$  time:
  - Minimum(), Maximum()
  - Successor(), Predecessor()
  - Search()
- Insert() and Delete():
  - Will also take  $O(\lg n)$  time
  - But will need special care since they modify tree

## Red-Black Trees: An Example

- *Color this tree:*




Red-black properties:

1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

## Red-Black Trees: The Problem With Insertion

- Insert 8

- *Where does it go?*




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

## Red-Black Trees: The Problem With Insertion

- Insert 8

- *Where does it go?*

- *What color should it be?*




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

## Red-Black Trees: The Problem With Insertion

- Insert 8

- *Where does it go?*
- *What color should it be?*




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

## Red-Black Trees: The Problem With Insertion

- Insert 11

- *Where does it go?*




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

## Red-Black Trees: The Problem With Insertion

- Insert 11

- *Where does it go?*
- *What color?*




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

## Red-Black Trees: The Problem With Insertion

- Insert 11

- *Where does it go?*
- *What color?*
  - Can't be red! (#3)



1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black


## Red-Black Trees: The Problem With Insertion

- Insert 11

- *Where does it go?*

- *What color?*

- Can't be red! (#3)
    - Can't be black! (#4)



1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black


## Red-Black Trees: The Problem With Insertion

- Insert 11

- *Where does it go?*

- *What color?*

- Solution:  
recolor the tree




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black



## Red-Black Trees: The Problem With Insertion

- Insert 10

■ *Where does it go?*




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

## Red-Black Trees: The Problem With Insertion

- Insert 10

■ *Where does it go?*


■ *What color?*




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

## Red-Black Trees: The Problem With Insertion

- Insert 10
  - *Where does it go?*
  - *What color?*
    - A: no color! Tree is too imbalanced
    - Must change tree structure to allow recoloring
  - Goal: restructure tree in  $O(\lg n)$  time




Copyright © The McGraw-Hill Companies, Inc. Permission required for reproduction or display.



**Figure 13.2** The rotation operations on a binary search tree. The operation  $\text{LEFT-ROTATE}(T, x)$  transforms the configuration of the two nodes on the left into the configuration on the right by changing a constant number of pointers. The configuration on the right can be transformed into the configuration on the left by the inverse operation  $\text{RIGHT-ROTATE}(T, y)$ . The letters  $\alpha$ ,  $\beta$ , and  $\gamma$  represent arbitrary subtrees. A rotation operation preserves the binary-search-tree property: the keys in  $\alpha$  precede  $\text{key}[x]$ , which precedes the keys in  $\beta$ , which precede  $\text{key}[y]$ , which precedes the keys in  $\gamma$ .


## RB Trees: Rotation

- Our basic operation for changing tree structure is called *rotation*:




- Does rotation preserve inorder key ordering?
- What would the code for **rightRotate()** actually do?

## RB Trees: Rotation




- Answer: A lot of pointer manipulation
  - $x$  keeps its left child
  - $y$  keeps its right child
  - $x$ 's right child becomes  $y$ 's left child
  - $x$ 's and  $y$ 's parents change
- What is the running time?



**Figure 13.3** An example of how the procedure  $\text{LEFT-ROTATE}(T, x)$  modifies a binary search tree. Inorder tree walks of the input tree and the modified tree produce the same listing of key values.


## Rotation Example

- Rotate left about 9:



## Rotation Example

- Rotate left about 9:



## Red-Black Trees: Insertion

- Insertion: the basic idea
  - Insert  $x$  into tree, color  $x$  red
  - Only r-b property 3 might be violated (if  $p[x]$  red)
    - If so, move violation up tree until a place is found where it can be fixed
  - Total time will be  $O(\lg n)$

```

rbInsert(x)
    treeInsert(x);
    x->color = RED;
    // Move violation of #3 up tree, maintaining #4 as invariant:
    while (x!=root && x->p->color == RED)
        if (x->p == x->p->p->left)
            y = x->p->p->right;
            if (y->color == RED)
                x->p->color = BLACK;
                y->color = BLACK;
                x->p->p->color = RED;
                x = x->p->p;
            } Case 1
            else // y->color == BLACK
                if (x == x->p->right)
                    x = x->p;
                    leftRotate(x);
                    x->p->color = BLACK;
                    x->p->p->color = RED;
                    rightRotate(x->p->p);
                } Case 2
                } Case 3
        else // x->p == x->p->p->right
            (same as above, but with
             "right" & "left" exchanged)

```

```


rbInsert(x)
    treeInsert(x);
    x->color = RED;
    // Move violation of #3 up tree, maintaining #4 as invariant:
    while (x!=root && x->p->color == RED)
        if (x->p == x->p->p->left)
            y = x->p->p->right;
            if (y->color == RED)
                x->p->color = BLACK;
                y->color = BLACK;
                x->p->p->color = RED;
                x = x->p->p;
            } Case 1: uncle is RED
            else // y->color == BLACK
                if (x == x->p->right)
                    x = x->p;
                    leftRotate(x);
                    x->p->color = BLACK;
                    x->p->p->color = RED;
                    rightRotate(x->p->p);
                } Case 2
                } Case 3
        else // x->p == x->p->p->right
            (same as above, but with
             "right" & "left" exchanged)

```

## RB Insert: Case 1

```
if (y->color == RED)
    x->p->color = BLACK;
    y->color = BLACK;
    x->p->p->color = RED;
    x = x->p->p;
```

- Case 1: “uncle” is red
- In figures below, all  $\Delta$ ’s are equal-black-height subtrees




*Change colors of some nodes, preserving #4: all downward paths have equal b.h.  
The while loop now continues with x’s grandparent as the new x*

## RB Insert: Case 1

```
if (y->color == RED)
    x->p->color = BLACK;
    y->color = BLACK;
    x->p->p->color = RED;
    x = x->p->p;
```

- Case 1: “uncle” is red
- In figures below, all  $\Delta$ ’s are equal-black-height subtrees




*Same action whether x is a left or a right child*

## RB Insert: Case 2

```
if (x == x->p->right)
    x = x->p;
    leftRotate(x);
// continue with case 3 code
```

- Case 2:
  - “Uncle” is black
  - Node  $x$  is a right child
- Transform to case 3 via a left-rotation




*Transform case 2 into case 3 ( $x$  is left child) with a left rotation  
This preserves property 4: all downward paths contain same number of black nodes*

## RB Insert: Case 3

```
x->p->color = BLACK;
x->p->p->color = RED;
rightRotate(x->p->p);
```

- Case 3:
  - “Uncle” is black
  - Node  $x$  is a left child
- Change colors; rotate right



*Perform some color changes and do a right rotation  
Again, preserves property 4: all downward paths contain same number of black nodes*



## RB Insert: Cases 4-6

---

- Cases 1-3 hold if  $x$ 's parent is a left child
- If  $x$ 's parent is a right child, cases 4-6 are symmetric (swap left for right)

## Red-Black Trees: Deletion

---

- And you thought insertion was tricky...
- We will not cover RB delete in class
  - You should read section 14.4 on your own
  - Read for the overall picture, not the details

## Red-Black Trees

---

- Red-black trees do what they do very well
- *What do you think is the worst thing about red-black trees?*
- A: coding them up

## Skip Lists

## Skip Lists

---

- A relatively recent data structure
  - “A probabilistic alternative to balanced trees”
  - A randomized algorithm with benefits of r-b trees
    - $O(\lg n)$  expected time for Search, Insert
    - $O(1)$  time for Min, Max, Succ, Pred
  - *Much* easier to code than r-b trees
  - Fast!

The End

---