

B-Trees

Introduction


- Similar to Red-Black tree or other balanced search trees.
- Different from other balanced search tree, nodes may have many children. Data is stored in a disk. So mainly used for Disk I/O.
- Provides index structures for large amount of data.

Introduction – contd.

- Up to now, all data that has been stored in the tree has been in memory.
- If data gets too big for main memory, what do we do?
- All data cannot be resident in the main memory.
- If we keep a pointer to the tree in main memory, we could bring in just the nodes that we need.
- For instance, to do an insert with a BST, if we need the left child, we do a disk access and retrieve the left child.
- If the left child is NIL, then we can do the insert, and store the child node on the disk.
- Not too good for a BST


Introduction – contd.

- The problem with BST: storing the data requires disk accesses, which is expensive, compared to execution of machine instructions.
- If we can reduce the number of disk accesses, then the procedures run faster.
- The only way to reduce the number of disk accesses is to increase the number of keys in a node.
- The BST allows only one key per leaf.
- **Very good and often used for Search Engines!**
 - (when collection size gets very big → the index does not fit in memory)



Disk Operations

- In a disk, data is organized in disk pages.
- To read data in a disk,
 - The disk rotates, and the R/W head moves to the page containing the data.
 - Then, the whole disk page is read.



From: Cormen, T., C. Leiserson, R. Rivest, and C. Stein, *Introduction to Algorithms*, MIT Press, 2001.

An example of B Tree






Figure 18.1 A B-tree whose keys are the consonants of English. An internal node x containing $n[x]$ keys has $n[x] + 1$ children. All leaves are at the same depth in the tree. The lightly shaded nodes are examined in a search for the letter R .

B-Trees: Nodes

Internal node



leaf node



B-Trees: Nodes – contd.

A B-Tree is a rooted tree (whose root is $\text{root}[T]$) having the following properties:

1. Every internal node x has the following fields:

$n[x]$	$\text{key}_1[x]$	$\text{key}_2[x]$	$\text{key}_3[x]$	$\text{key}_4[x]$	$\text{key}_5[x]$	$\text{key}_6[x]$	$\text{key}_7[x]$	$\text{key}_8[x]$	$\text{leaf}[x]$
$c_1[x]$	$c_2[x]$	$c_3[x]$	$c_4[x]$	$c_5[x]$	$c_6[x]$	$c_7[x]$	$c_8[x]$	$c_9[x]$	

$n[x]$ is the number of keys in the node. $n[x] = 8$ above.

$\text{leaf}[x] = \text{false}$ for internal nodes, since x is not a leaf.

The $\text{key}_i[x]$ are the values of the keys, where $\text{key}_i[x] \leq \text{key}_{i+1}[x]$.

$c_i[x]$ are pointers to child nodes. All the keys in $c_i[x]$ have values that are between $\text{key}_{i-1}[x]$ and $\text{key}_i[x]$.

B-Trees: Nodes – contd.

☀ Leaf nodes have no child pointers


☀ $\text{leaf}[x] = \text{true}$ for leaf nodes.

☀ All leaf nodes are at the same level

$n[x]$	$\text{key}_1[x]$	$\text{key}_2[x]$	$\text{key}_3[x]$	$\text{key}_4[x]$	$\text{key}_5[x]$	$\text{key}_6[x]$	$\text{key}_7[x]$	$\text{key}_8[x]$	$\text{leaf}[x]$

Properties of B-Trees - contd.

For any node n in a B-tree T



- $n.key_i \leq c.key_1 \leq c.key_2 \leq \dots \leq c.key_q \leq n.key_{i+1}$
- If n is a leaf node, the depth of n is h , where h is the tree's height.

Properties of B-Trees contd.

- Root must have 1 key .
- The **minimum degree**, t , of the B-tree is the lower bound on the number of keys a node can contain. ($t \geq 2$)
- Every node other than the root must have at least $t - 1$ keys. Every internal node other than the root thus has at least t children.
- If the tree is nonempty, the root must have at least one key.
- Every node can contain at most $2t - 1$ keys. Therefore, an internal node can have at most $2t$ children.
- We say that a node is **full** if it contains exactly $2t - 1$ keys.




Figure 18.3 A B-tree of height 2 containing over one billion keys. Each internal node and leaf contains 1000 keys. There are 1001 nodes at depth 1 and over one million leaves at depth 2. Shown inside each node x is $n[x]$, the number of keys in x .





Figure 18.4 A B-tree of height 3 containing a minimum possible number of keys. Shown inside each node x is $n[x]$.

🔦 Thm :

If $n \geq 1$, then for any n -key B-tree T of height h and minimum degree $t \geq 2$,

$$h \leq \log_t \frac{n+1}{2}.$$

Proof :



$$\begin{aligned} n &\geq 1 + (t-1) \sum_{i=1}^h 2t^{i-1} \\ &= 1 + 2(t-1) \cdot \left(\frac{t^h - 1}{t - 1} \right) = 2t^h - 1. \\ \frac{n+1}{2} &\geq t^h. \quad \log_t \frac{n+1}{2} \geq h. \end{aligned}$$

Height of B-Tree

- 🔦 If $n \geq 1$, then for any n -key B-tree of height h and minimum degree $t \geq 2$,

$$\text{height} = h \leq \log_t [(n+1)/2]$$

- 🔦 The important thing to notice is that the height of the tree is \log base t . So, as t increases, the height, for any number of nodes n , will decrease.

- 🔦 Using the formula $\log_a x = (\log_b x) / (\log_b a)$, we can see that

$$\blacksquare \log_2 10^6 = (\log_{10} 10^6) / (\log_{10} 2) \approx 6 / 0.30102999566398 \approx 19$$

$$\blacksquare \log_{10} 10^6 = 6$$

So, 13 less disk accesses to get to the leafs!

B-tree of degree t

- Insertion and deletions:
 - More complicate but still $\log(n)$
 - Split and merge operation.

-
- Convention :
 - Root of the B-tree is always in main memory.
 - Any nodes that are passed as parameters must already have had a DISK_READ operation performed on them.
 - Operations :
 - Searching a B-Tree.
 - Creating an empty B-tree.
 - Splitting a node in a B-tree.
 - Inserting a key into a B-tree.
 - Deleting a key from a B-tree.

🍌 B-Tree-Search(x,k) :

📦 Algorithm :

```
B-Tree-Search(x,k)
{  i ← 1
  while i ≤ n[x] and k > keyi[x]
    do i ← i + 1
  if i ≤ n[x] and k = keyi[x]
    then return(x,i)
  if leaf[x] then return NULL
    else DISK - READ(Ci[x])
    return B - Tree - Search(Ci[x],k)
}
```

$O(th) = O(t \log_t n)$.

🍌 B-Tree-Created(T) :

📦 Algorithm :


```
B-Tree-Create(T)
{  x ← Allocate - Node()
  Leaf[x] ← TRUE
  n[x] ← 0
  DISK - WRITE(x)
  root[T] ← x
}
```

$O(1)$

📦 time :

● B-Tree-Split-Child(x,i,y) :

■ Splitting a node in a B-Tree :



Splitting a full node y (have $2t-1$ keys) around its median key $key_t[y]$ into 2 nodes having $(t-1)$ keys each.

■ Algorithm :


```

B-Tree-Split-Child(x,i,y)
{
    z ← Allocate-Node()
    leaf[z] ← leaf[y]
    n[z] ← t - 1
    for j ← 1 to t-1 do key_j[z] ← key_{j+t}[y]
    if not leaf[y] then
        for j ← 1 to t do C_j[z] ← C_{j+t}[y]
    n[y] ← t - 1
    for j ← n[x] + 1 downto i + 1 do C_{j+1}[x] ← C_j[x]
    C_{j+1}[x] ← z
    for j ← n[x] downto i do Key_{j+1}[x] ← Key_j[x]
    Key_i[x] ← Key_i[y]
    n[x] ← n[x] + 1
    DISK - WRITE(y)
    DISK - WRITE(z)
    DISK - WRITE(x)
}

```

● B-Tree-Insert(T, k) :

✚ Insert a key in a B-Tree :



✚ Algorithm :

```

B-Tree-Insert( $T, k$ )
{
     $r \leftarrow \text{root}[T]$ 
    if  $n[r] = 2t - 1$  then
    {
         $S \leftarrow \text{Allocate-Node}()$ 
         $\text{root}[T] \leftarrow S$ 
         $\text{leaf}[S] \leftarrow \text{FALSE}$ 
         $n[S] \leftarrow 0$ 
         $C_i[S] \leftarrow r$ 
        B-Tree-Split-Child( $S, i, r$ )
        B-Tree-Insert-Nonfull( $S, k$ )
    }
    else B-Tree-Insert-Nonfull( $r, k$ )
}
  
```

B-Tree-Insert-Nonfull(x,k) :

Algorithm :

B-Tree-Insert-Nonfull(x,k)


```

{   i ← n[x]
    if leaf[x] then
        { while i ≥ 1 and k < keyi[x]
          do { keyi+1[x] ← keyi[x]
              i ← i - 1 }
          keyi+1[x] ← k
          n[x] ← n[x] + 1
          DISK - WRITE(x) }
    else
        { while i ≥ 1 and k < keyi[x]
          do i ← i - 1
          i ← i + 1
          DISK - READ(Ci[x])
          if n[Ci[x]] = 2t - 1
              then B - Tree - Split - Child(x,i,Ci[x])
                  if k > keyi[x] then i ← i + 1
                  B-Tree-Insert-Nonfull(Ci[x],k) }
    }
```


Example : Inserting keys into a B-Tree.

t=3



(a) Initial tree




(b) B inserted




(c) Q inserted



(d) L insert




(e) F insert



Deleting a key from a B-Tree :



(x has $\geq t$ keys)

1. K is in x and x is a leaf :




2. K is in x and x is an internal node :


a. **Recursively delete k' and replace k by k' in x .**




c. if both y, z has $t-1$ keys.




3. If K is not in internal node x :



a. If $C_i[x]$ has only $t-1$ keys but has an immediate sibling with



- Move a key from x down to $C_i[x]$.
- Move a key from $C_i[x]$'s immediate sibling to x .
- Move an appropriate child to $C_i[x]$ from its immediate sibling.




b. If $C_i[x]$ and all of $C_i[x]$'s immediate siblings have $t-1$ keys, merge $C_i[x]$ with one immediate sibling.


Example : Deleting a key from a B-Tree.

$t=3$

(a) Initial tree



(b) F delete : case 1



(c) M delete : case 2a

