

CS 540

Computer Networks II

Sandy Wang
sandy.w@svuca.edu

5. TUNNELS

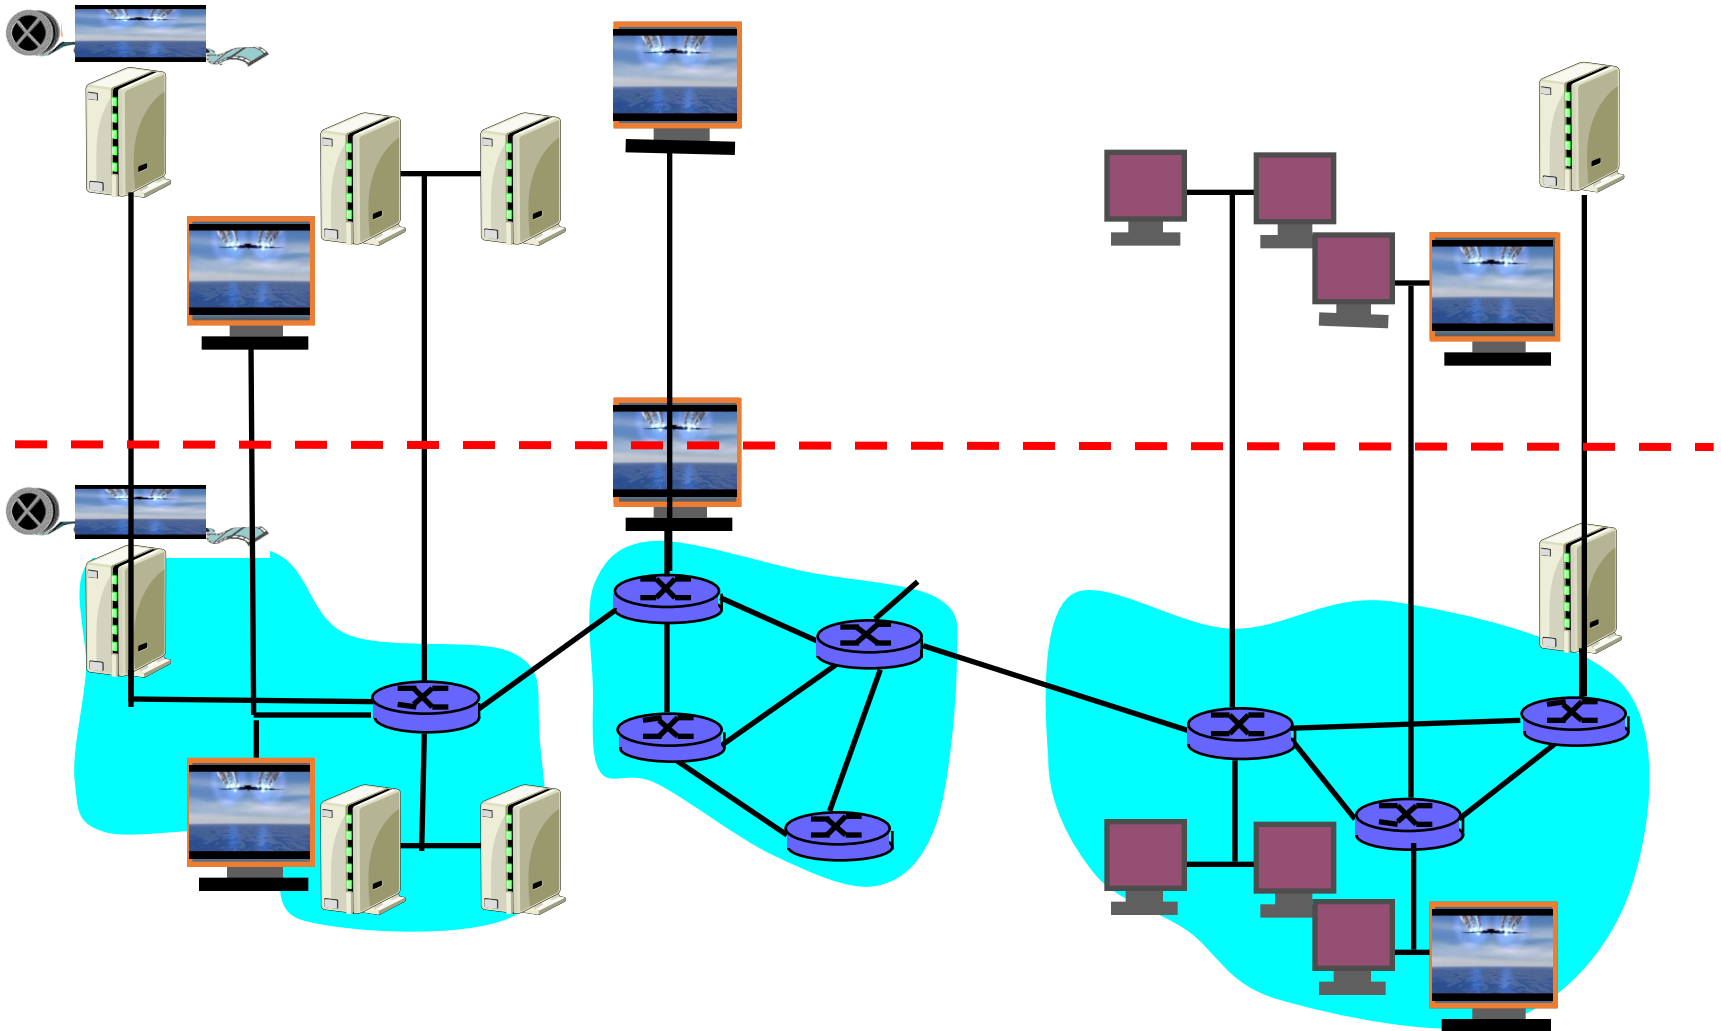
Topics

1. Overview
2. LAN Switching
3. IPv4
4. IPv6
- 5. Tunnels**
6. Routing Protocols -- RIP, RIPng
7. Routing Protocols -- OSPF
8. IS-IS
9. Midterm Exam
10. BGP
11. MPLS
12. Transport Layer -- TCP/UDP
13. Congestion Control & Quality of Service (QoS)
14. Access Control List (ACL)
15. Final Exam

Reference Books

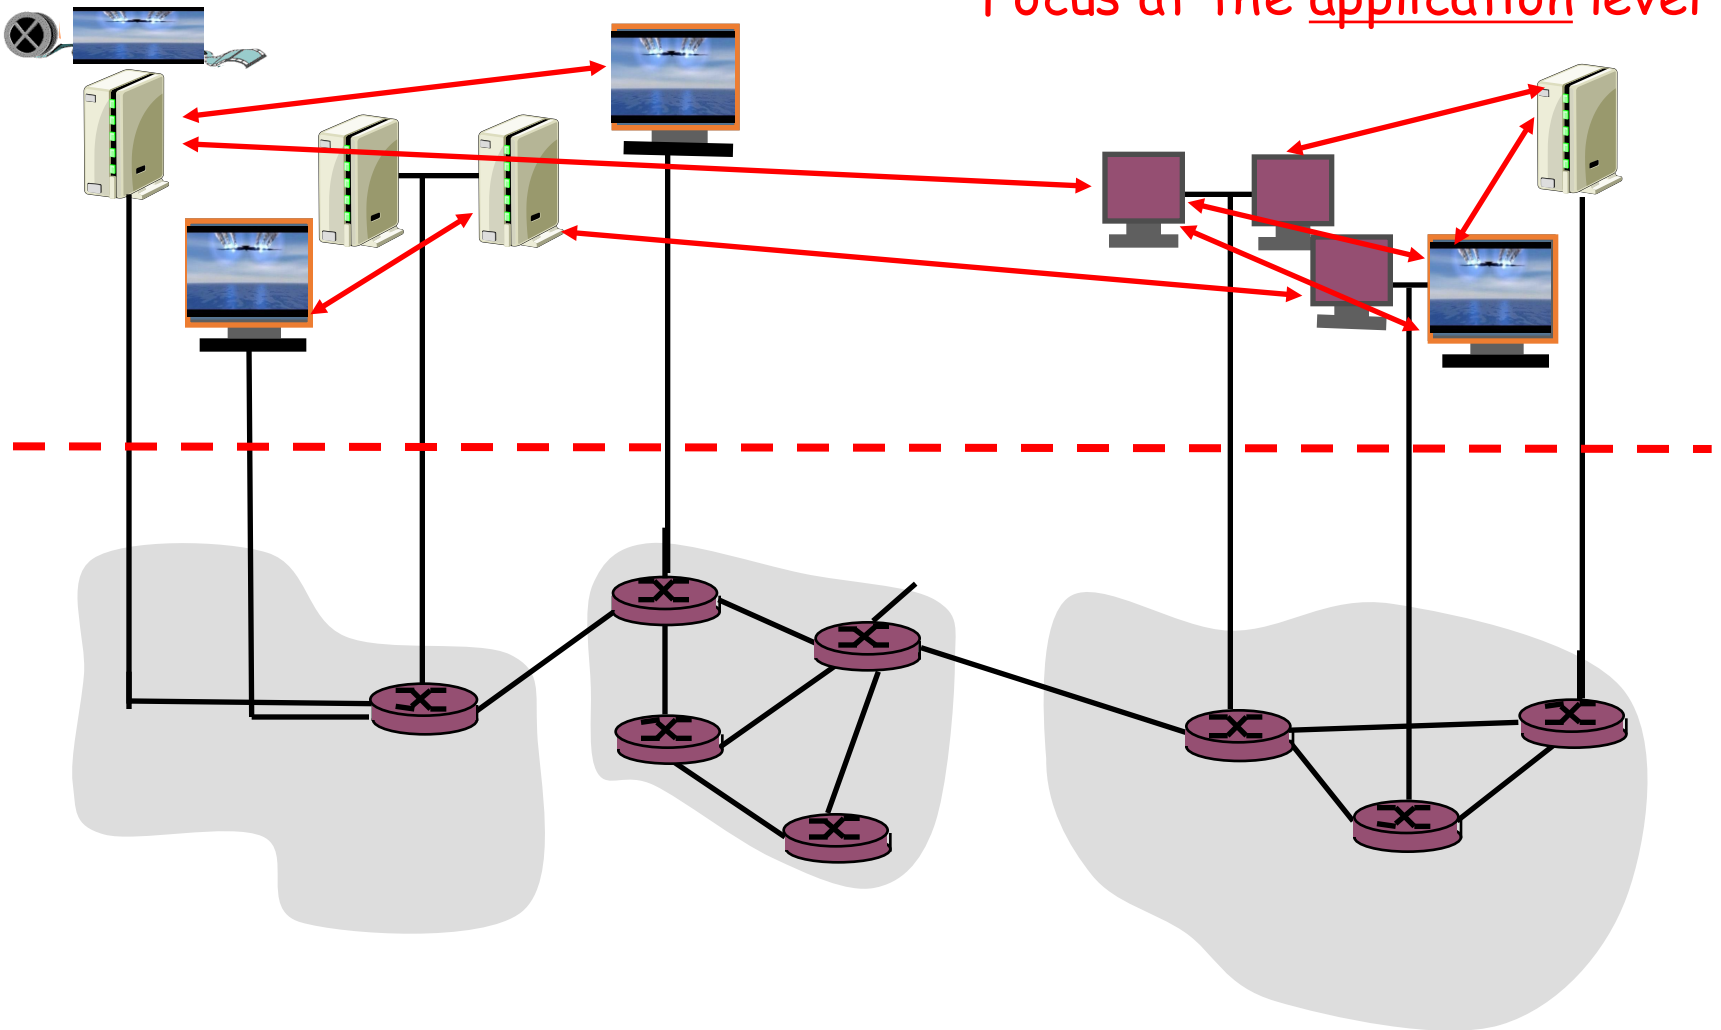
- **Routing TCP/IP Volume I, 2nd Edition** by Jeff Doyle and Jennifer Carroll
ISBN: 1-57870-089-2
- **Routing TCP/IP Volume II** by Jeff Doyle and Jennifer DeHaven
ISBN: 1-57870-089-2
- **Cisco CCNA Routing and Switching ICND2 200-101 Official Cert Guide, Academic Edition** by Wendel Odom -- July 10, 2013.
ISBN-13: 978-1587144882
- **The TCP/IP Guide: A Comprehensive, Illustrated Internet Protocols Reference** by Charles M. Kozierok – October 1, 2005.
ISBN-13: 978-1593270476
- **CCNA Routing and Switching 200-120 Network Simulator.** By Wendell Odom, Sean Wilkins. Published by Pearson IT Certification.
- <http://class.svuca.edu/~sandy/class/CS540/>

Overlay Networks



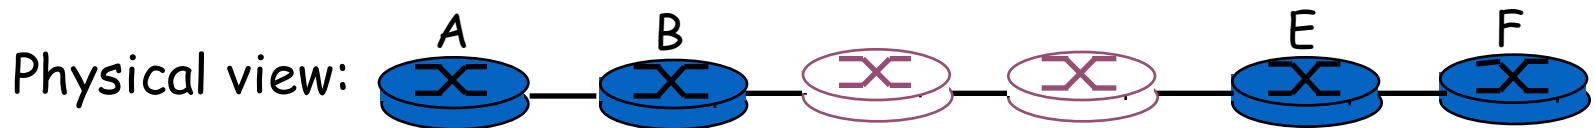
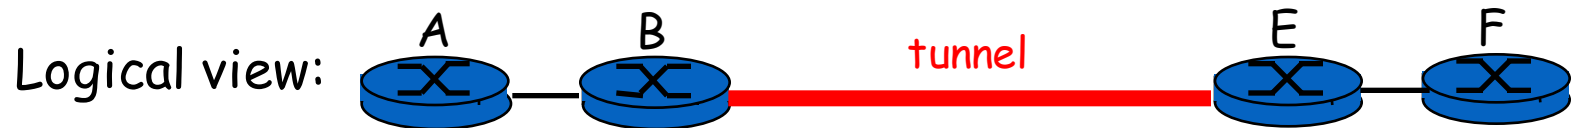
Overlay Networks

Focus at the application level



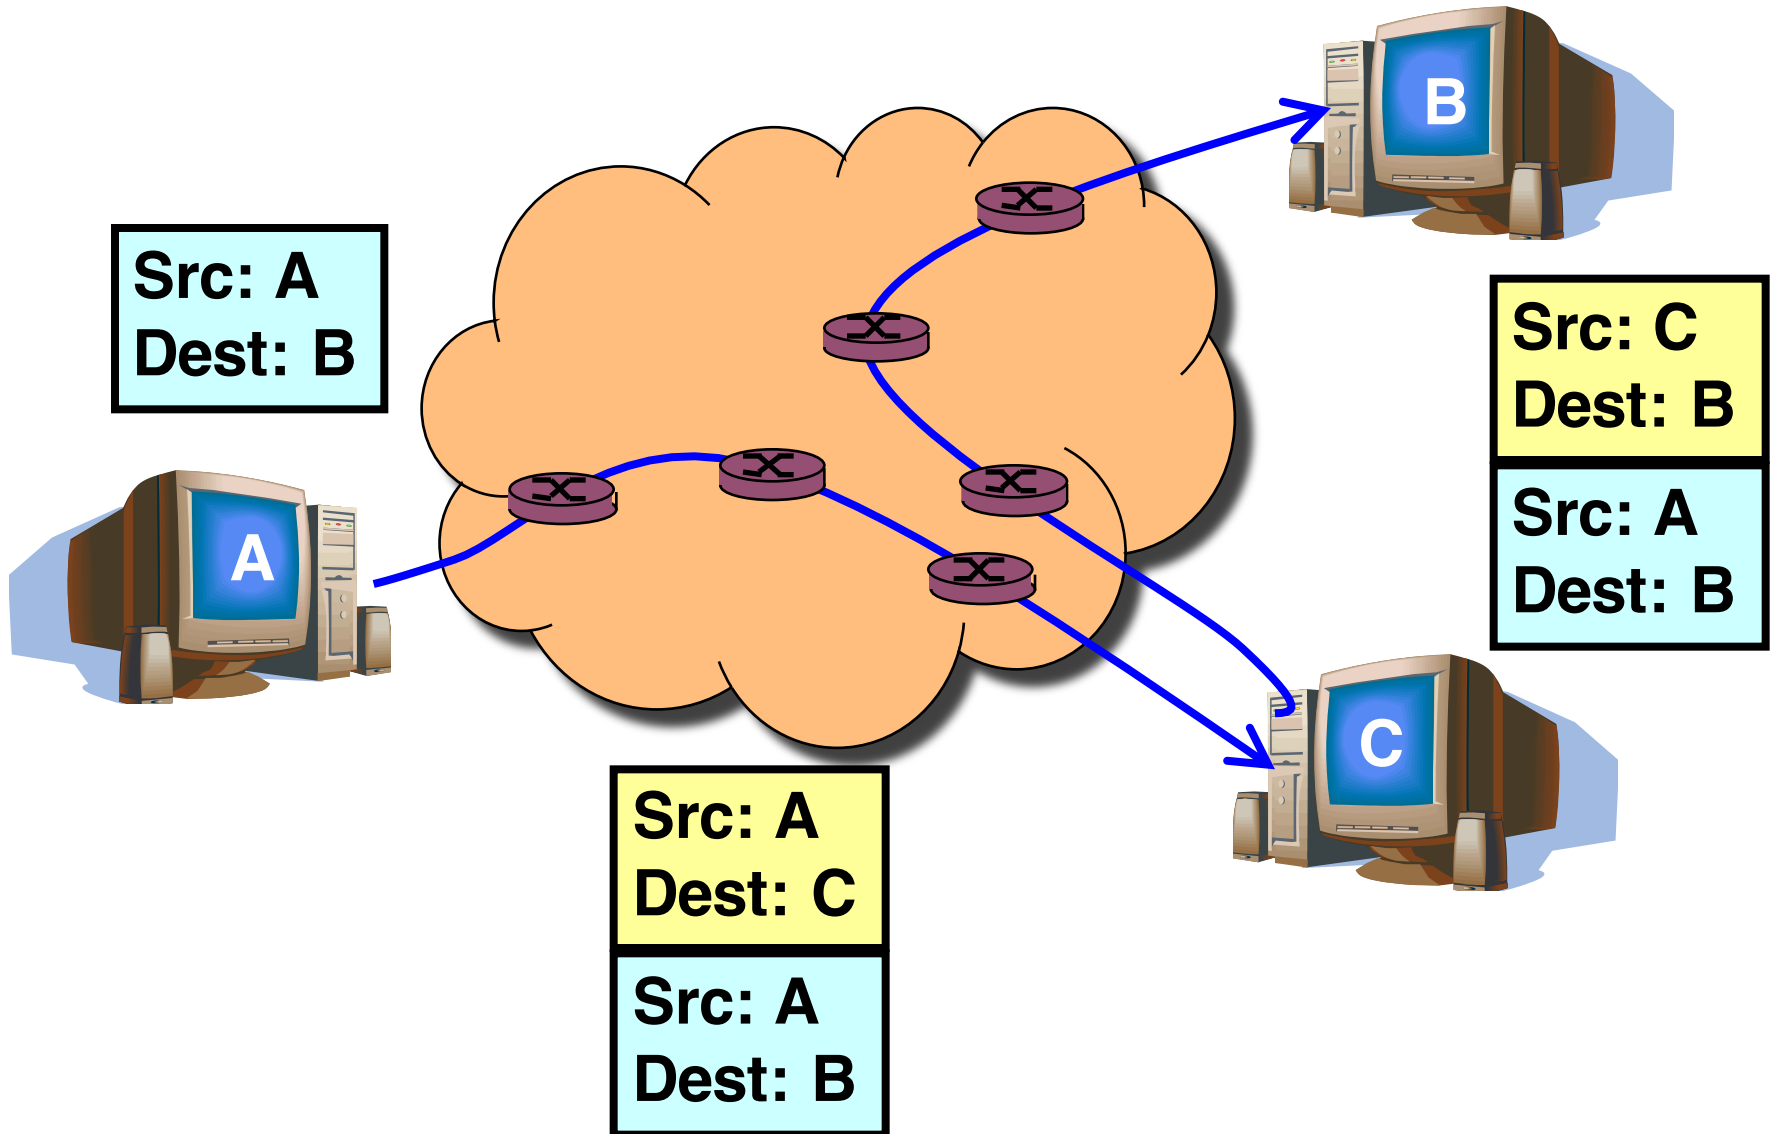
IP Tunneling to Build Overlay Links

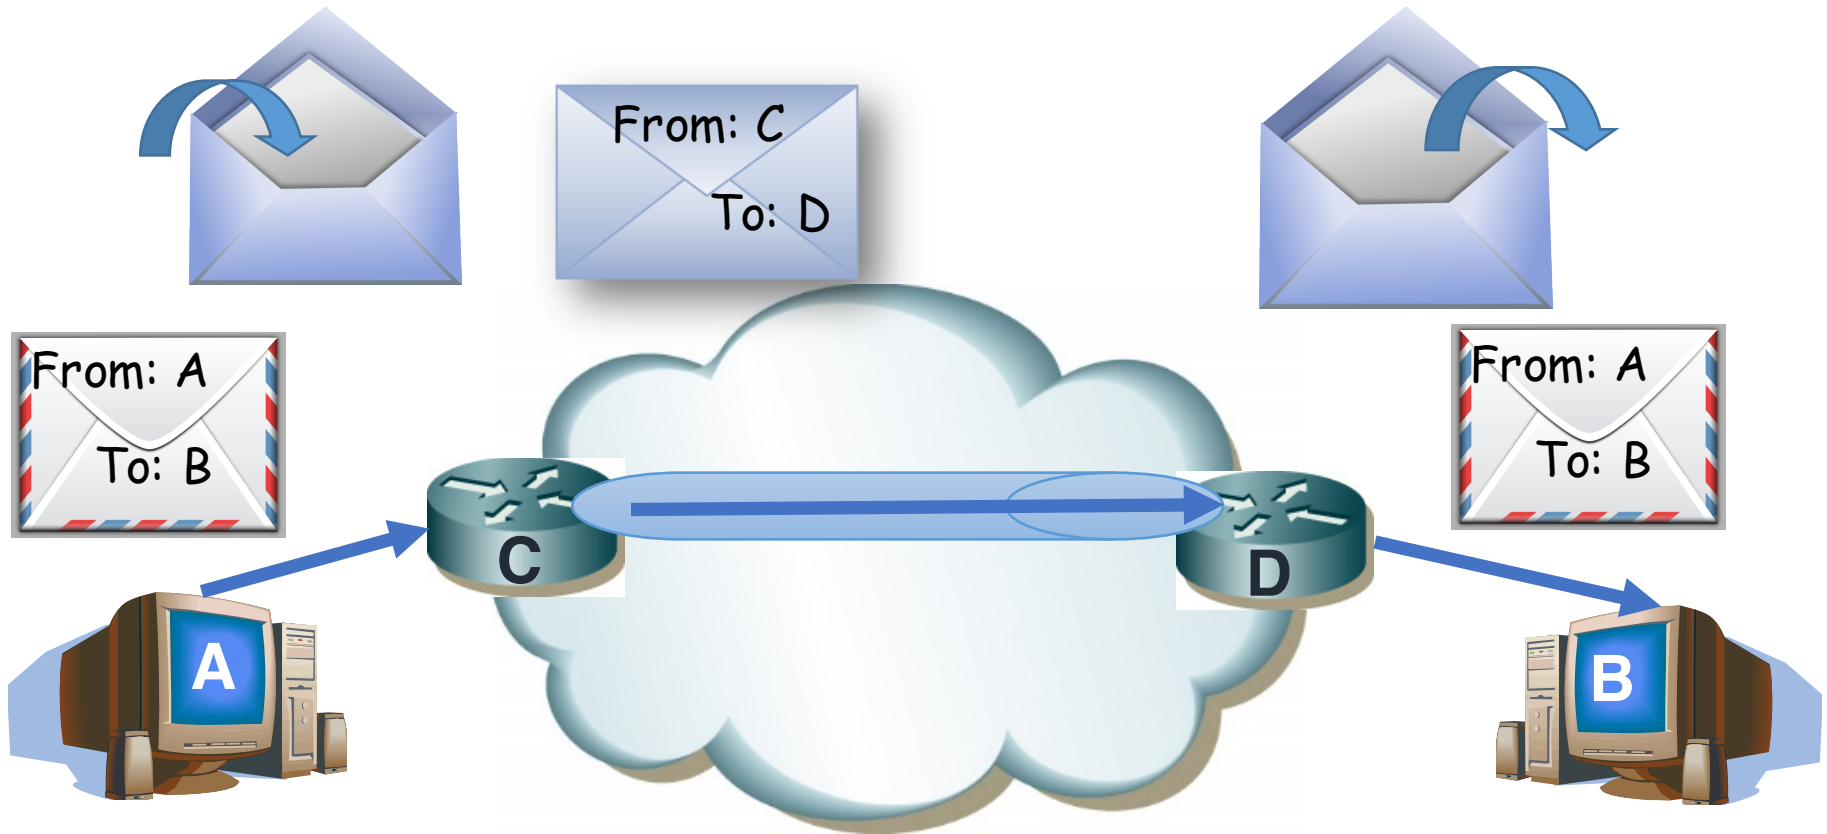
- IP tunnel is a virtual point-to-point link
 - Illusion of a direct link between two separated nodes



- Encapsulation of the packet inside an IP datagram
 - Node B sends a packet to node E
 - ... containing another packet as the payload

Tunnels Between End Hosts





Overlay Networks

- A logical network built on top of a physical network
 - Overlay links are tunnels through the underlying network
- Many logical networks may coexist at once
 - Over the same underlying network
 - And providing its own particular service

Common Uses

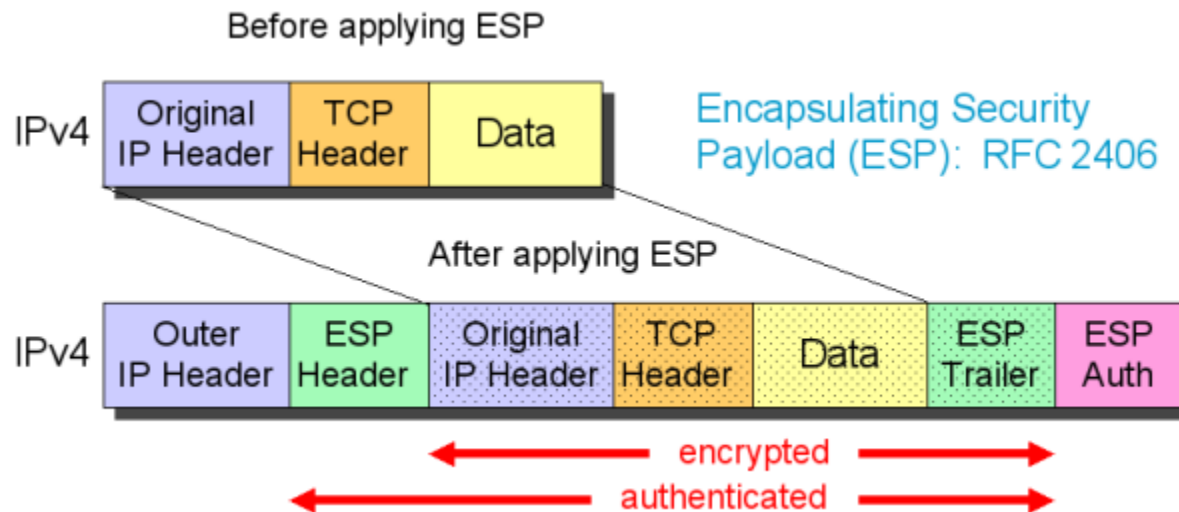
- Carry data over incompatible delivery-networks
- Provide a (encrypted) path through a public network
- Allowing “some kind” of traffic may lead to “any kind”

Misuse of Network Tunneling

- Pre-existing network-based security tools (firewalls, IDS) may not be able to apply the controls to the tunneled traffic
 - Evading traffic regulation
- Lack of host-based security controls
 - Defense in depth
- Inability for ingress and egress filtering
- ‘Open-ended’ tunnel may forward traffic to other internal hosts

IPSec Tunnel Mode

- The original IP packet is encrypted
- The ESP header indicates that the entire packet is the payload (IP-in-IP)
- Inserts a new IP header (next header is ESP)

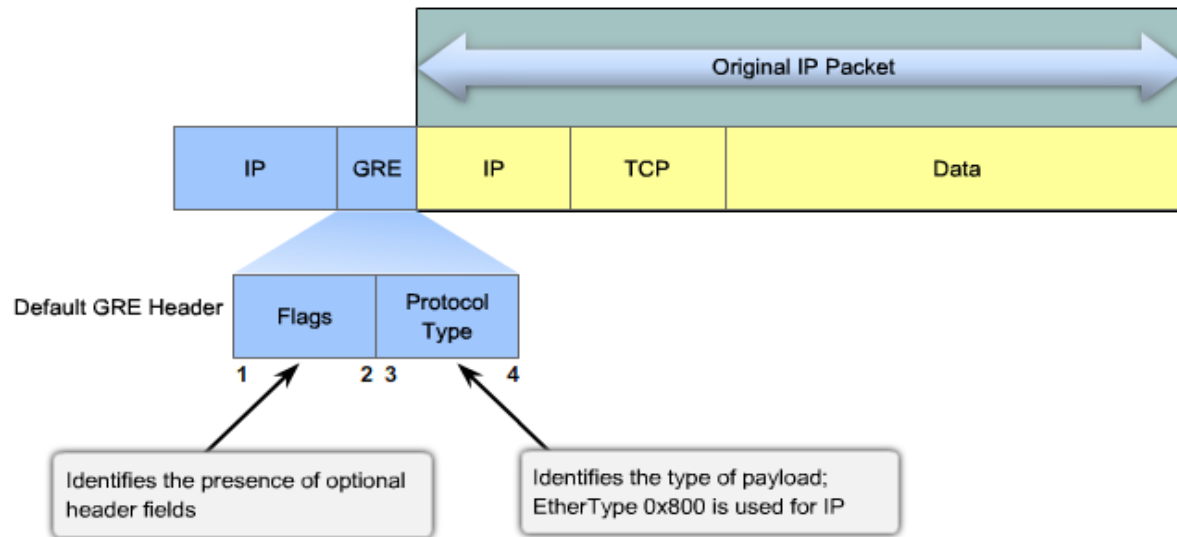


IPSec Tunnel Mode

- Security services from gateway to gateway or from host to gateway over an insecure network
- The entire original packet is encrypted
 - Internal traffic behind the gateways is not protected
- Often used to implement Virtual Private Networks (IPsec VPNs)
 - Site-to-site
 - Client-to-site

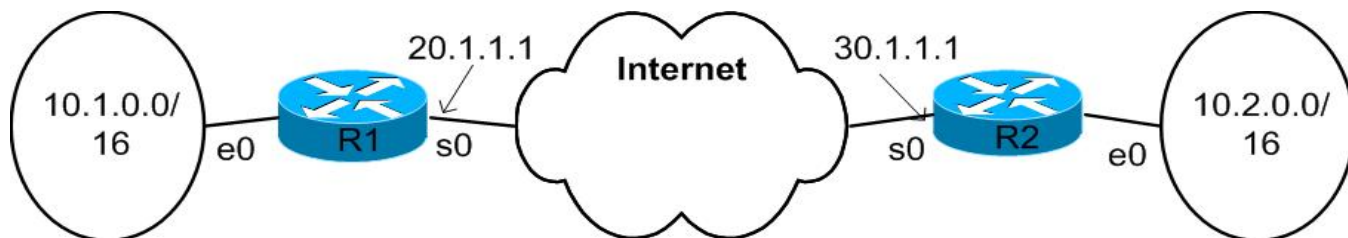
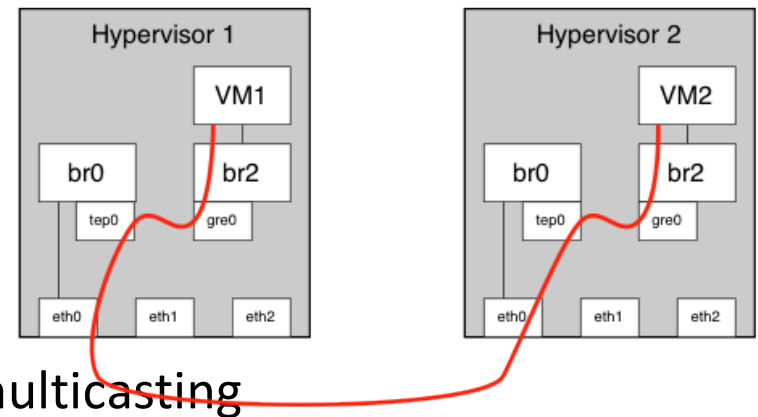
GRE – Generic Routing Encapsulation

- “GRE (Generic Routing Encapsulation) specifies a protocol for encapsulation of an arbitrary protocol over another arbitrary network layer protocol” – RFC 2784 and 2890
- Point-to-point links



GRE and IP

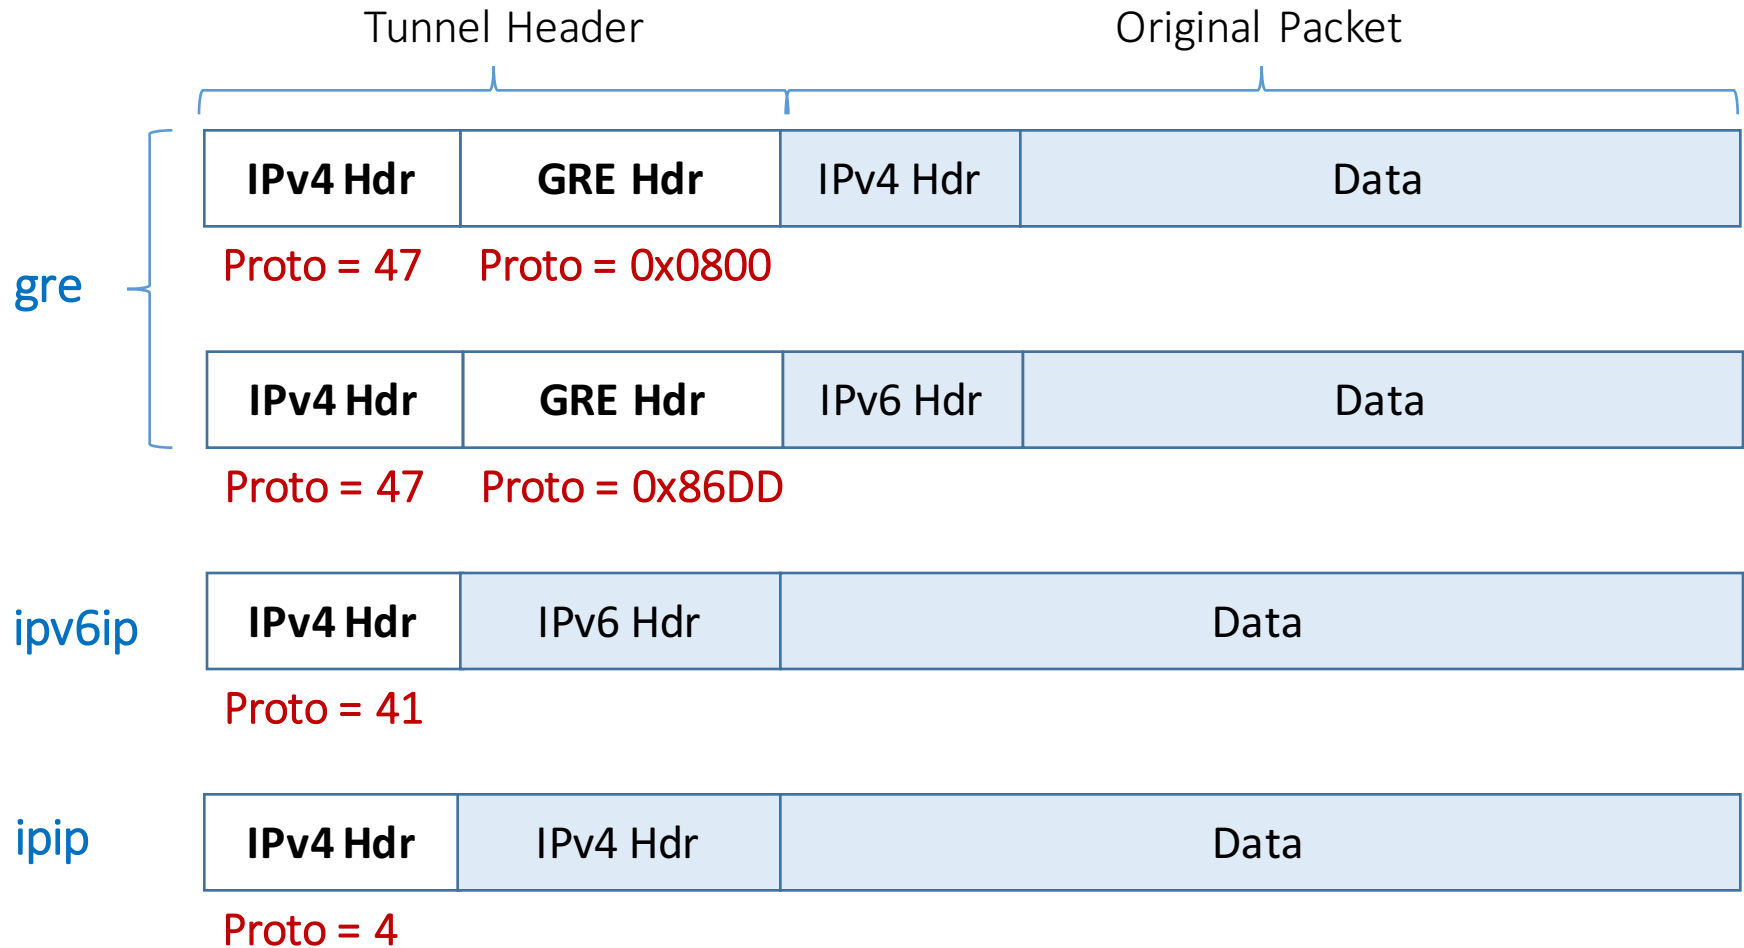
- Ethernet over IPv4/IPv6
(e.g. Openstack Neutron)
- Support for tunneling broadcasting/multicasting
 - e.g. Delivering routing updates to multiple sites
- IPv4/IPv6 over IPv4/IPv6
- No default encryption/security services
 - IPSec Tunnel/Transport over GRE



IP Protocol Numbers

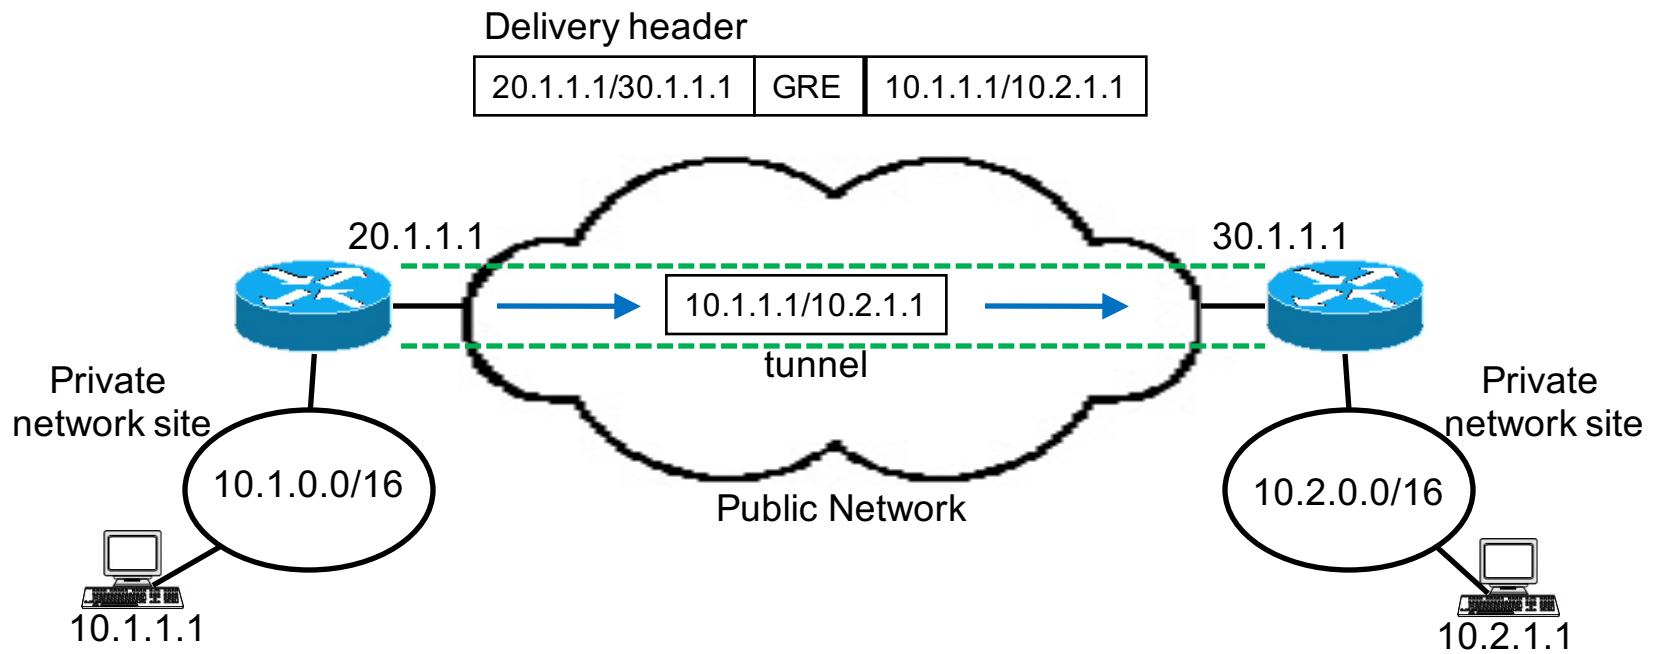
| Decimal | Hex | Keyword | Protocol |
|---------|------|----------|---|
| 4 | 0x04 | IP-in-IP | <u>IP in IP</u> (encapsulation) |
| 41 | 0x29 | IPv6 | IPv6 Encapsulation |
| 47 | 0x2F | GRE | <u>Generic Routing Encapsulation</u> |
| 50 | 0x32 | ESP | <u>Encapsulating Security Payload</u> |
| 51 | 0x33 | AH | <u>Authentication Header</u> |

Tunnel Mode and Encapsulation



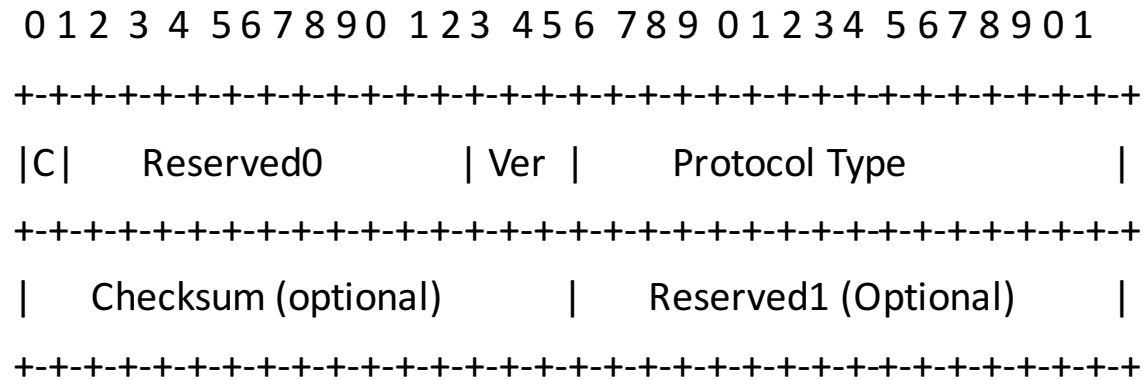
Generic Routing Encapsulation (GRE)

- Tunneling
 - Encapsulation with delivery header
 - The addresses in the delivery header are the addresses of the head-end and the tail-end of the tunnel



Generic Routing Encapsulation (GRE)

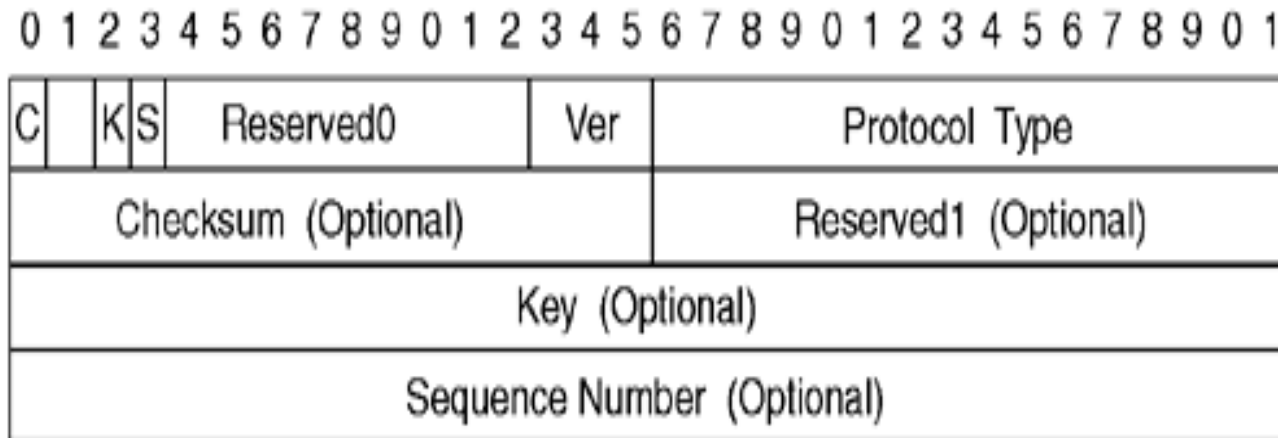
- RFC 2784 – Generic Routing Encapsulation



Protocol Type – EtherType in RFC 1700. Examples below:

| EtherType | Protocol |
|-----------|--|
| 0x0800 | Internet Protocol version 4 (IPv4) |
| 0x0806 | Address Resolution Protocol (ARP) |
| 0x86DD | Internet Protocol Version 6 (IPv6) |

RFC 2890 – Key and Sequence Number Extensions to GRE

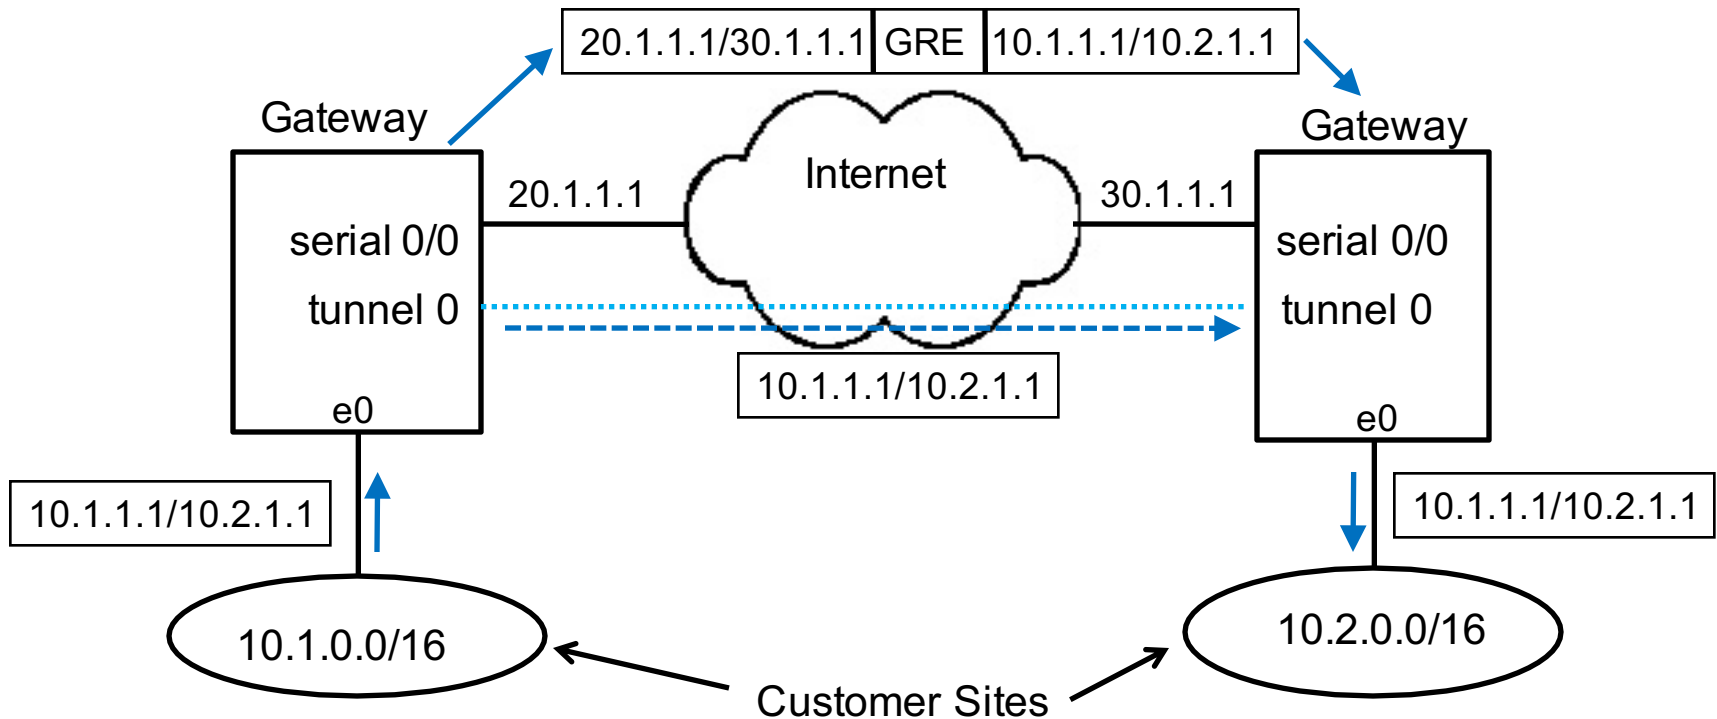


Key – Identify an individual traffic flow within a tunnel. Packets belonging to a traffic flow are encapsulated using the same Key value and the decapsulating tunnel endpoint identifies packets belonging to a traffic flow based on the Key Field value.

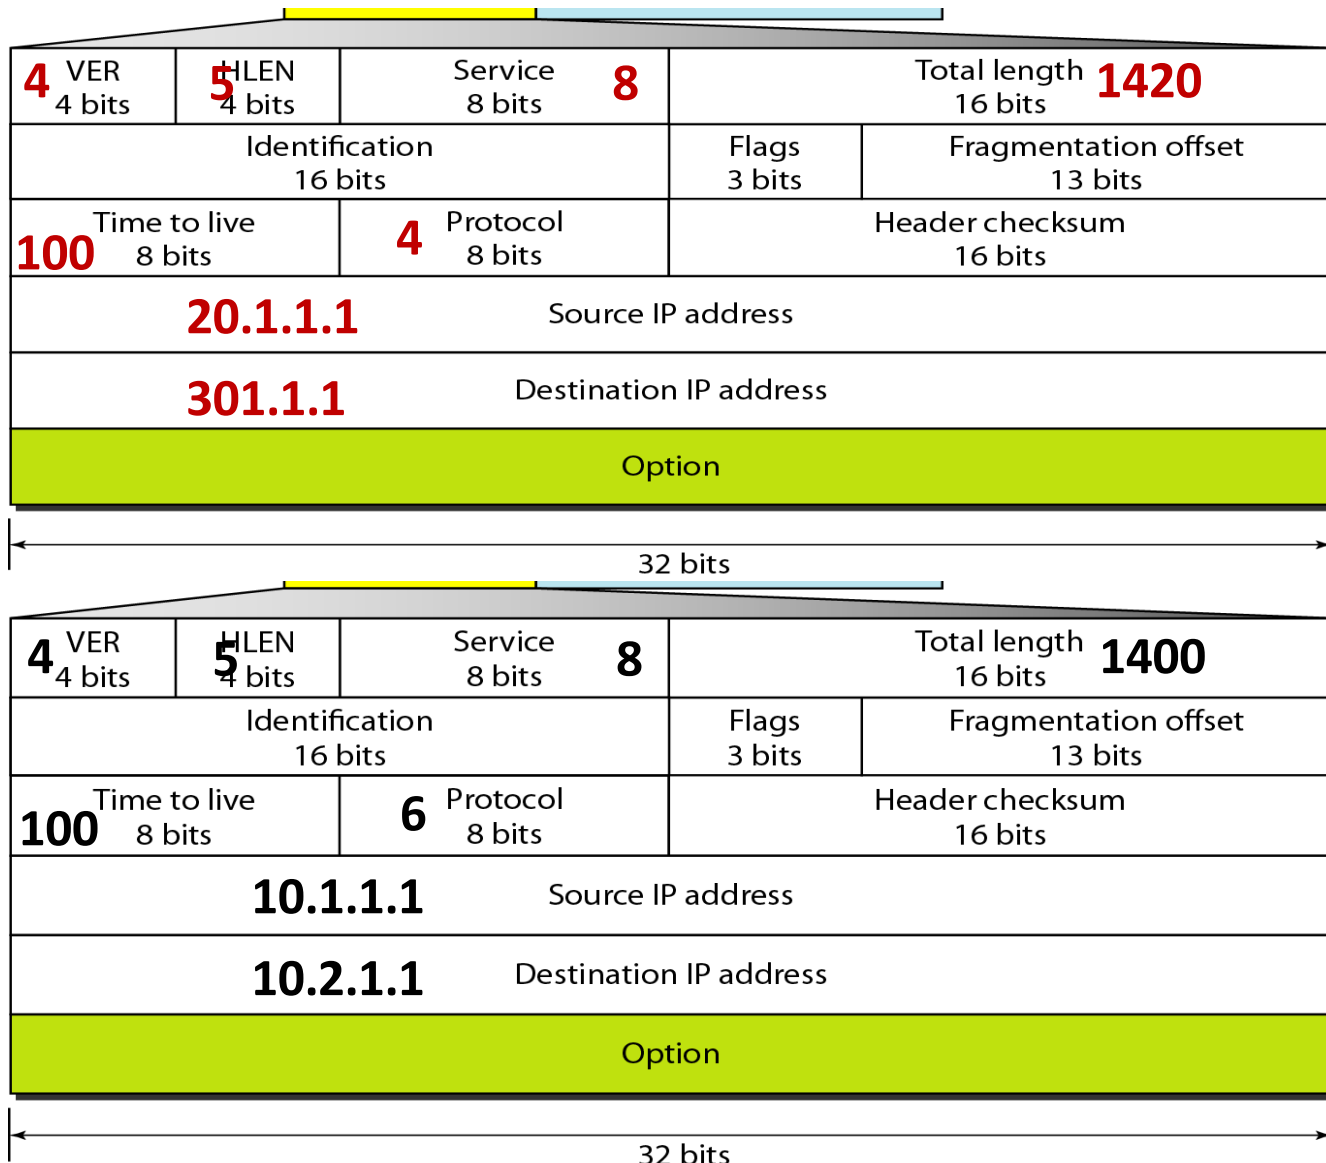
Sequence Number -- inserted by the encapsulator when Sequence Number Present Bit is set. The Sequence Number **MUST** be used by the receiver to establish the order

Generic Routing Encapsulation (GRE)

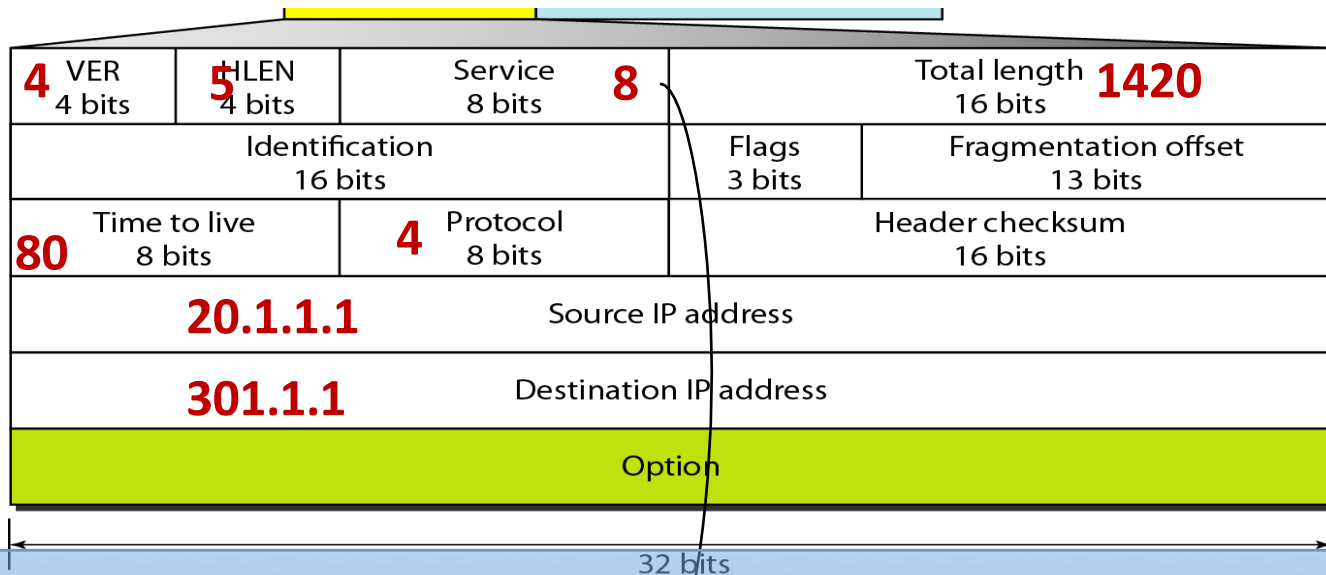
- IP access of the tunnel through the tunnel interface



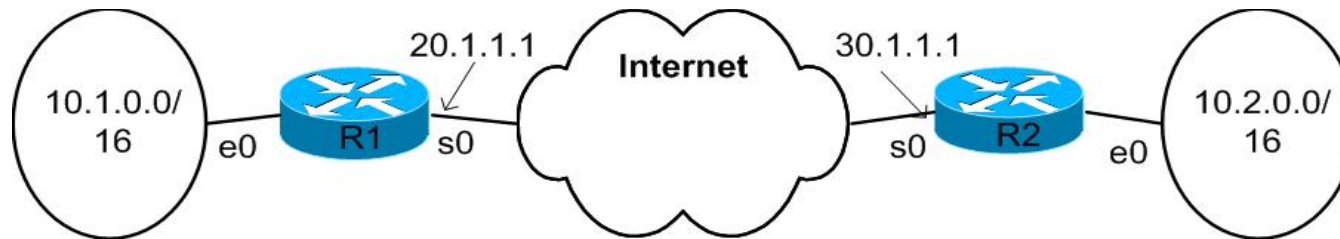
Encapsulation



Decapsulation



Generic Routing Encapsulation (GRE)



```
interface tunnel0
mode GRE
tunnel source 20.1.1.1
tunnel destination 30.1.1.1
!  
ip route 10.2.0.0 255.255.0.0 tunnel0
```

```
interface tunnel0
mode GRE
tunnel source 30.1.1.1
tunnel destination 20.1.1.1
!  
ip route 10.1.0.0 255.255.0.0 tunnel0
```

Routing table of R1

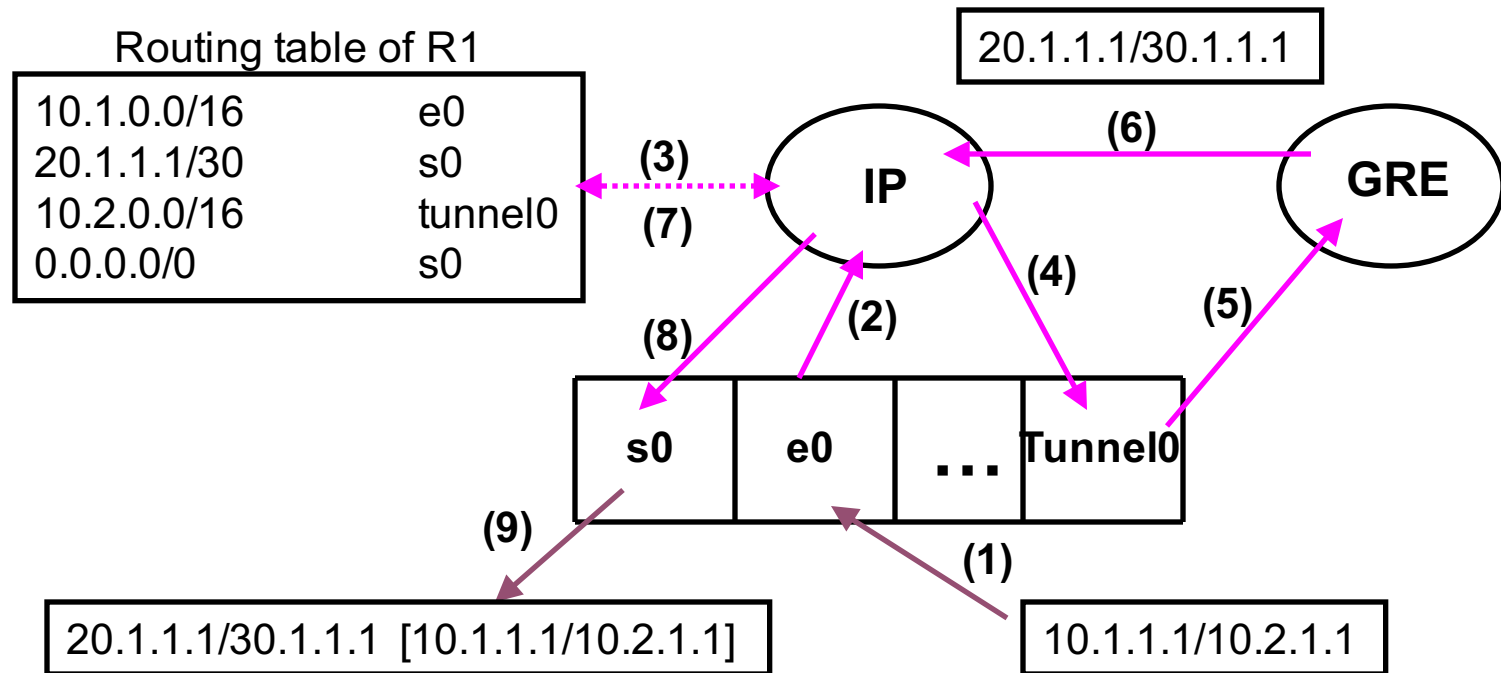
| | |
|-------------|---------|
| 10.1.0.0/16 | e0 |
| 20.1.1.1/30 | s0 |
| 10.2.0.0/16 | tunnel0 |
| 0.0.0.0/0 | s0 |

Routing table of R2

| | |
|-------------|---------|
| 10.2.0.0/16 | e0 |
| 30.1.1.1/30 | s0 |
| 10.1.0.0/16 | tunnel0 |
| 0.0.0.0/0 | s0 |

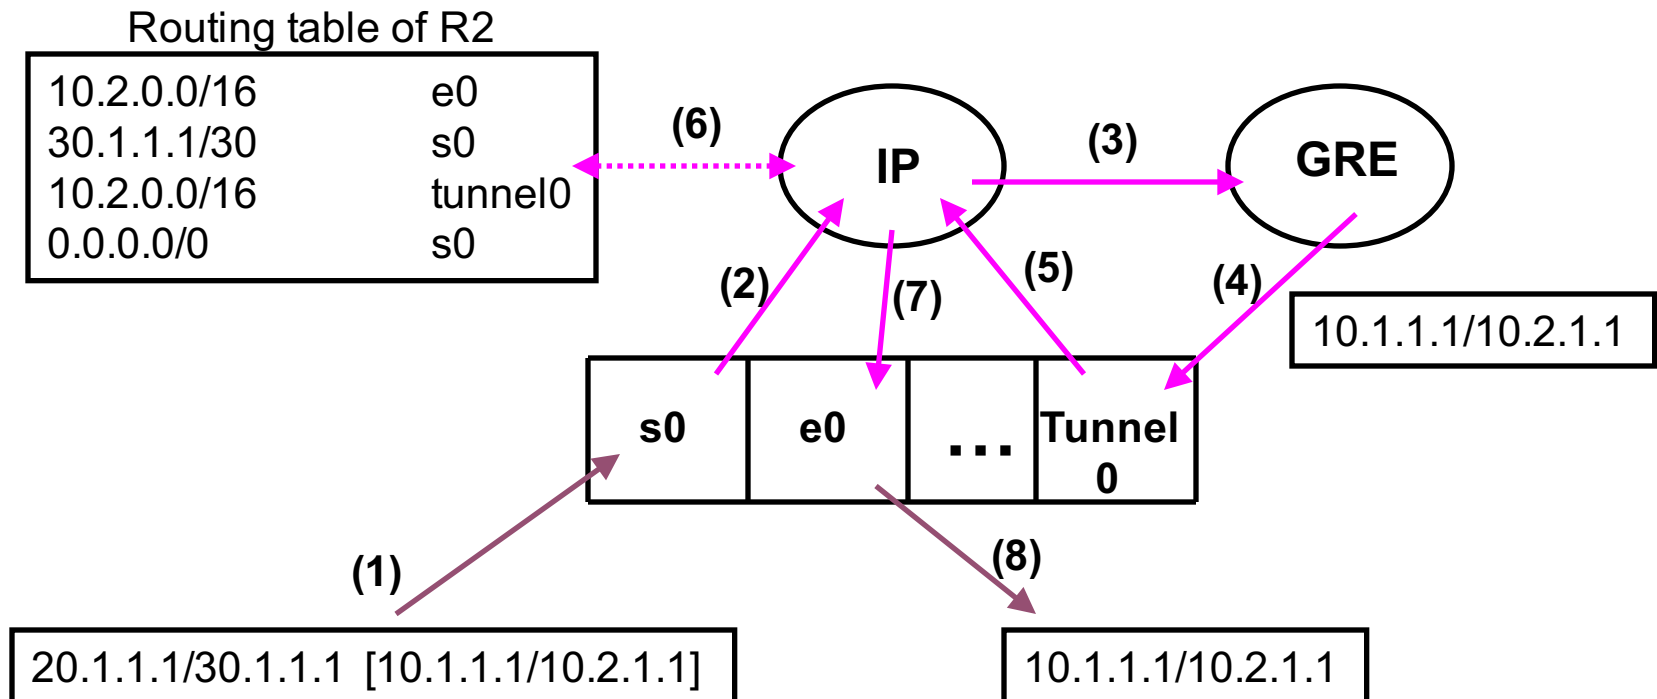
Generic Routing Encapsulation (GRE)

- Tunneling mechanism at IP
 - Outbound traffic



Generic Routing Encapsulation (GRE)

- Inbound traffic



IP Tunnel Example -- Linux

```
ip [ OPTIONS ] tunnel { COMMAND | help }
```

```
ip tunnel { add | change | del | show | prl } [ NAME ]
```

```
    [ mode MODE ] [ remote ADDR ] [ local ADDR ]
```

```
    [ [i|o]seq ] [ [i|o]key KEY ] [ [i|o]csum ] ]
```

```
    [ encaps limit ELIM ] [ ttl TTL ]
```

```
    [ tos TOS ] [ flowlabel FLOWLABEL ]
```

```
    [ prl-default ADDR ] [ prl-nodetault ADDR ] [ prl-delete ADDR]
```

```
    [ [no]pmtudisc ] [ dev PHYS_DEV ]
```

IP Tunnel Example -- Linux

MODE := { ipip | gre | sit | isatap | ip6ip6 | ipip6 | ip6gre | any }

ADDR := { IP_ADDRESS | any }

TOS := { STRING | 00..ff | inherit | inherit/STRING | inherit/00..ff }

ELIM := { none | 0..255 }

TTL := { 1..255 | inherit }

KEY := { DOTTED_QUAD | NUMBER }

TIME := NUMBER[s|ms]

Transformation From IPv4 to IPv6

Three strategies have been devised by the IETF to provide for a smooth transition from IPv4 to IPv6.

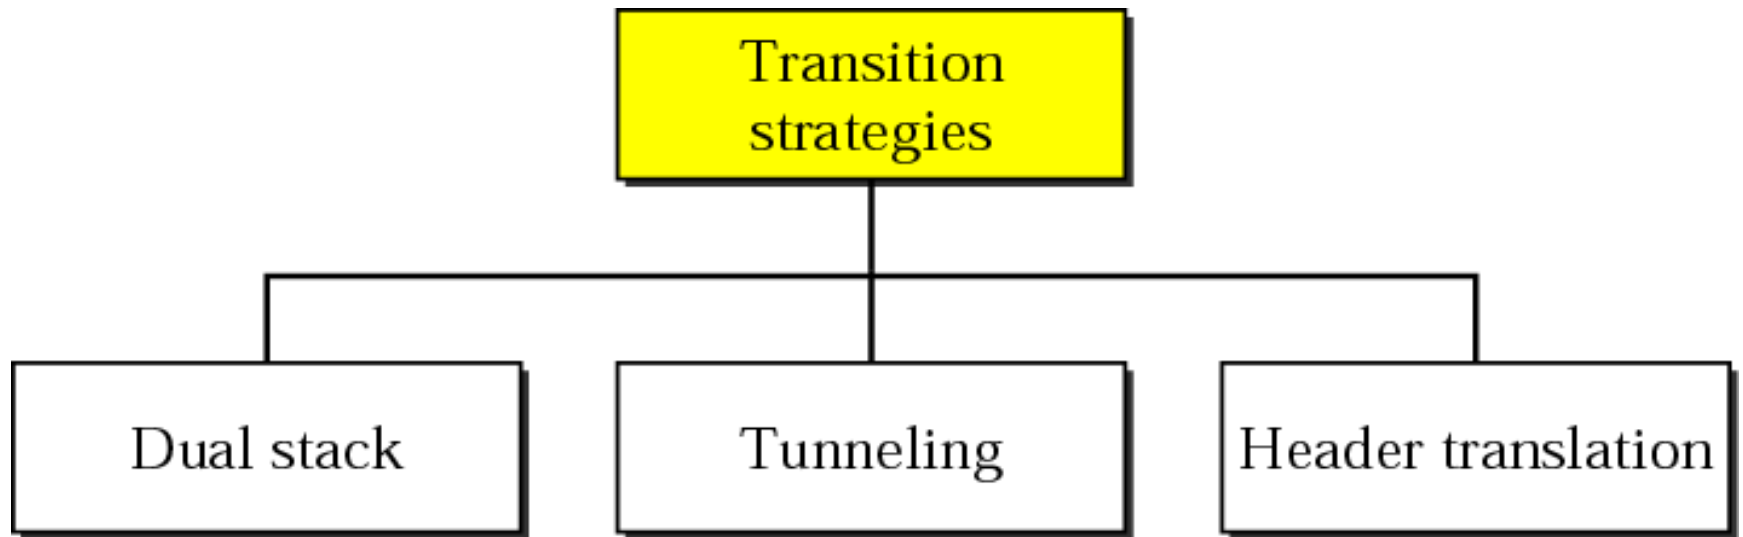
The topics discussed in this section include:

Dual Stack

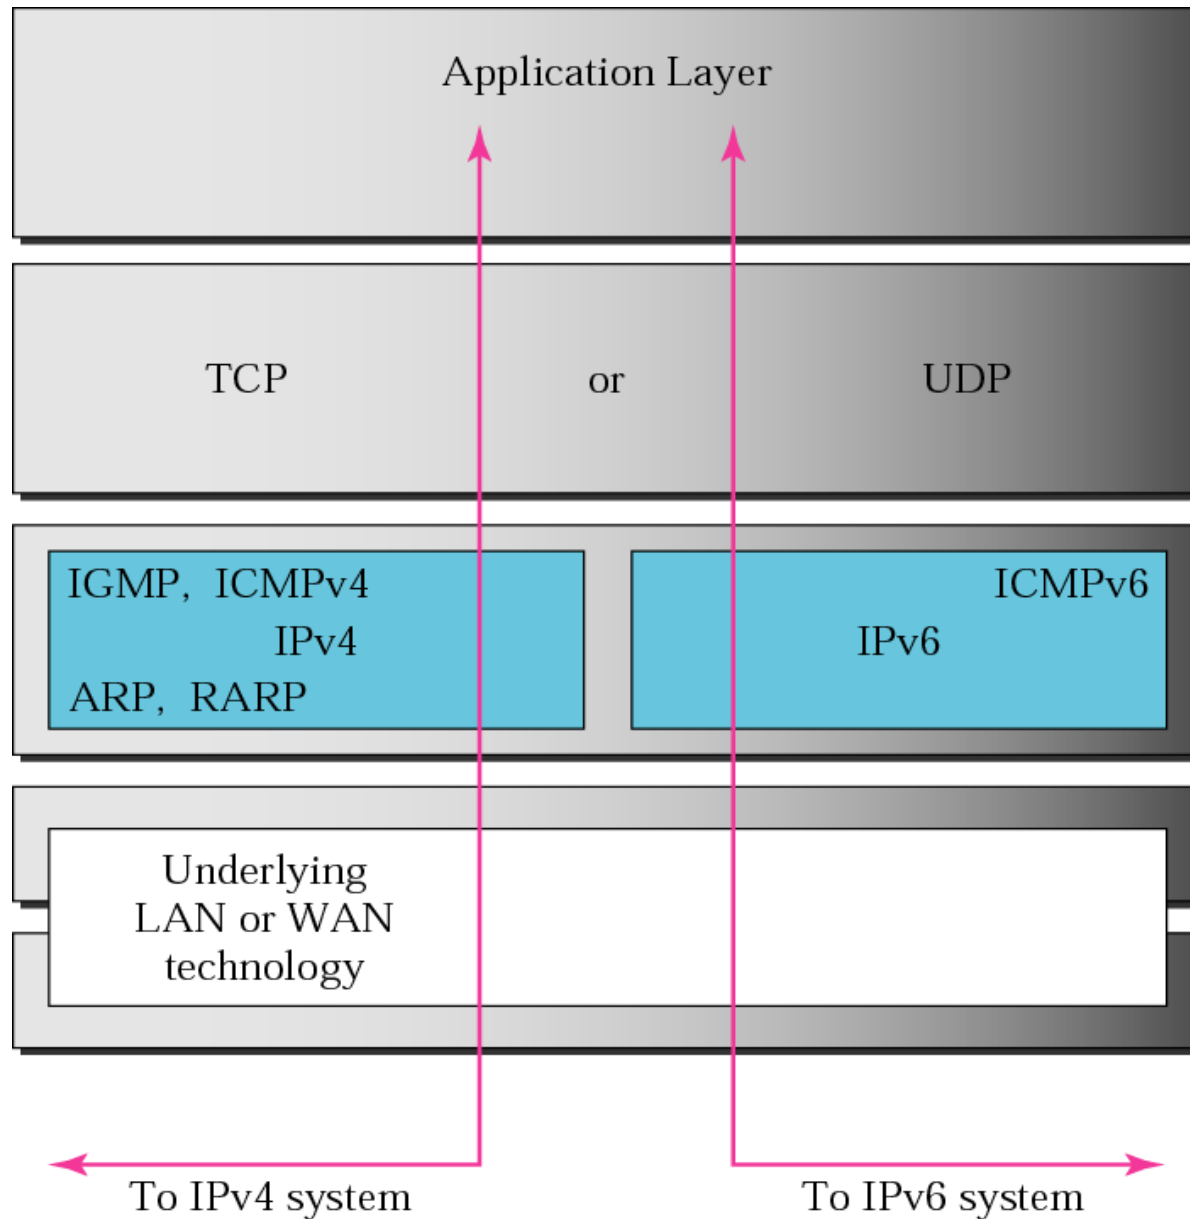
Tunneling

Header Translation

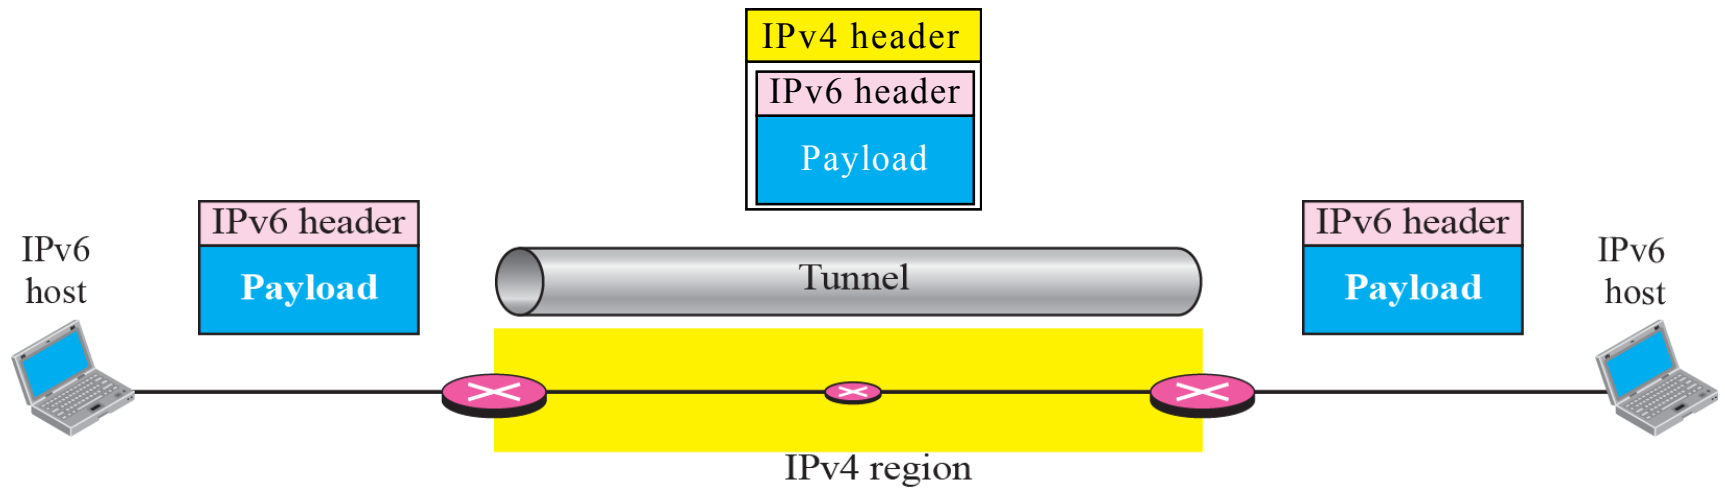
Three transition strategies



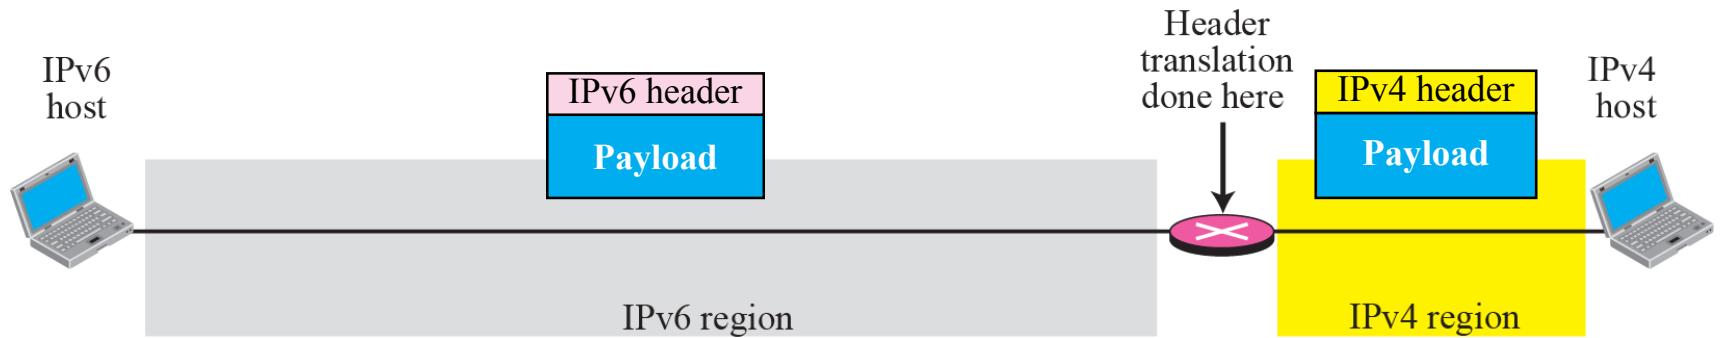
Dual stack



Tunneling strategy



Header translation strategy

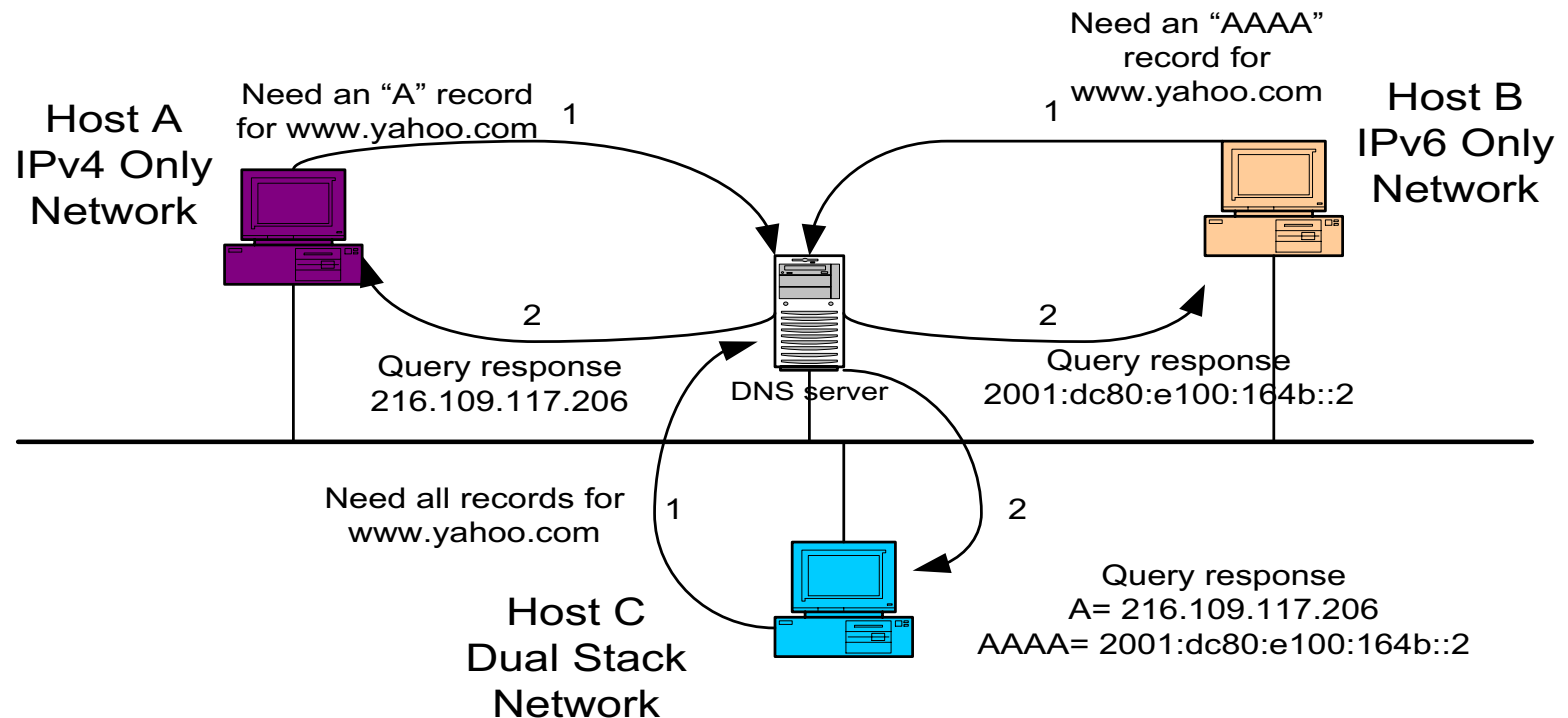


Naming Services

- DNS must be included in transition strategy
- Resolving Names:
 - IPv4 specifies “A” records
 - IPv6 specifies “AAAA” records
- Applications should be aware of both records
- Will require development update and thorough testing
- Tools like “Scrubber” by Sun make it easy

Naming Services

Querying DNS server



IPv6 over IPv4 Transition Mechanisms

- Tunnel Brokers provide a network tunneling service
- 6in4 – IPv6 over IPv4

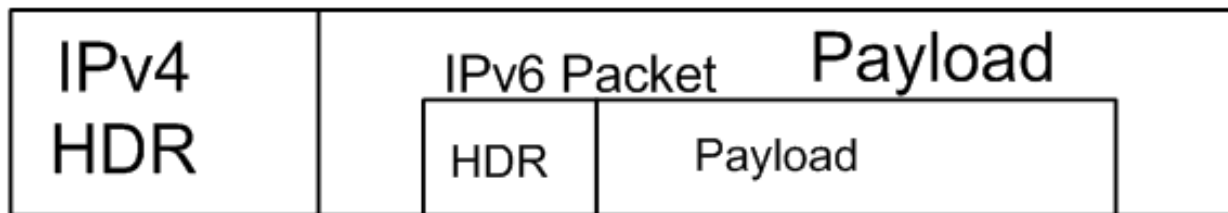


- Configured Tunnel
- 6to4 Tunnel
- ISATAP
- 6RD Tunnel
- Teredo – IPv6 over UDP over IPv4
- ...and others

Manually Configured Tunnels

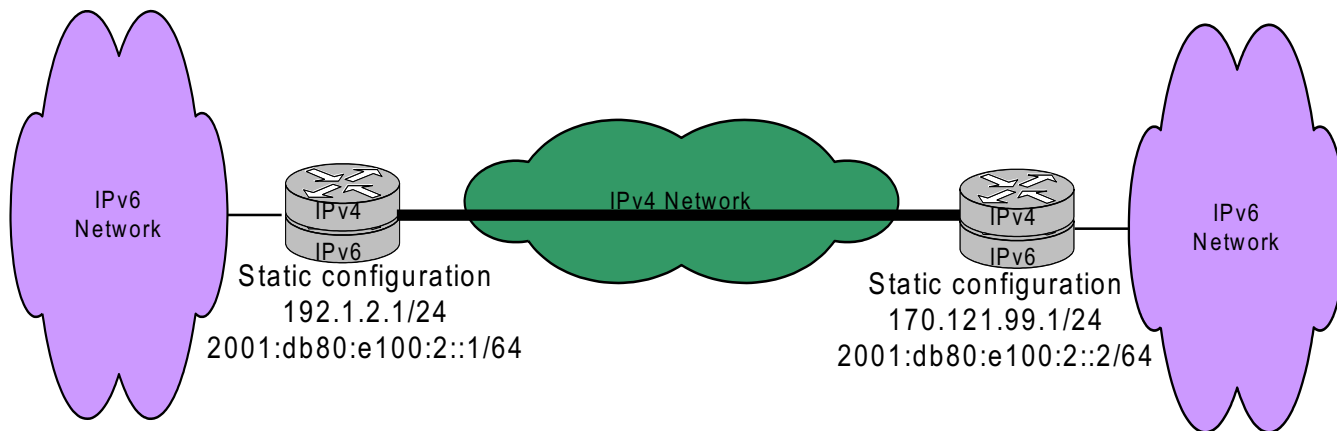
- Manually configured tunnels are logical tunnels formed when one protocol version packet is encapsulated in the payload of another version packet
- e.g. IPv4 encapsulated in IPv6 or IPv6 encapsulated in IPv4

IPv4 Packet with tunneling

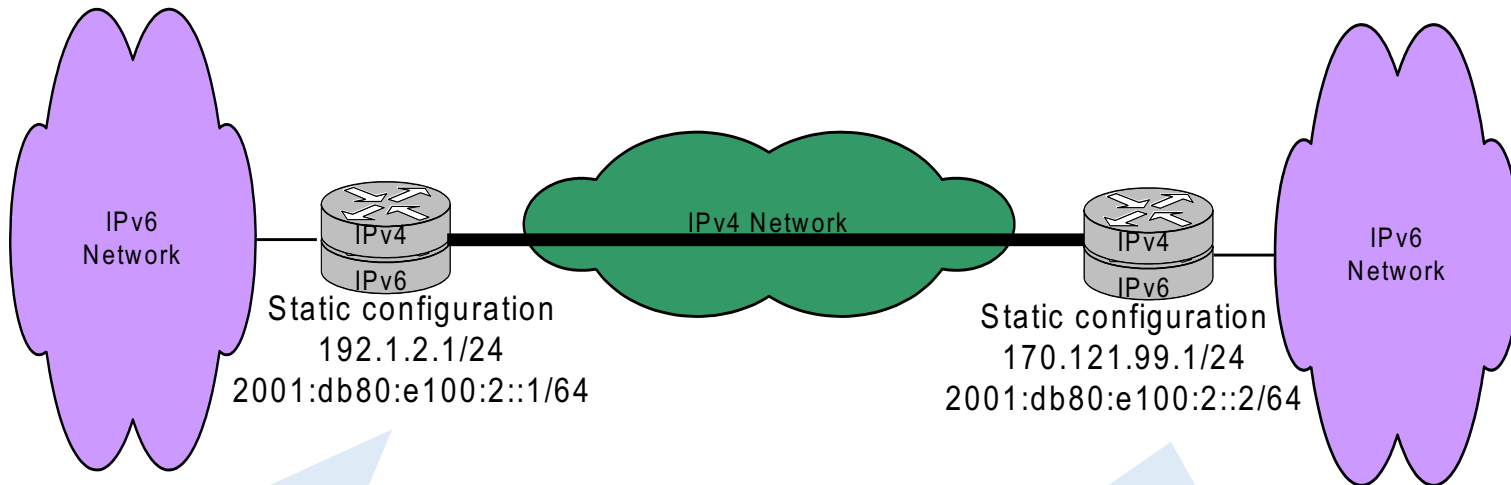


Configured Tunnel-building

- Configured tunnels require static IPv4 addresses
- Configured tunnels are generally setup and maintained by a network administrator
- Configured tunnels are a proven IPv6 deployment technique and provide stable links



Configured Tunnel -- Configuration

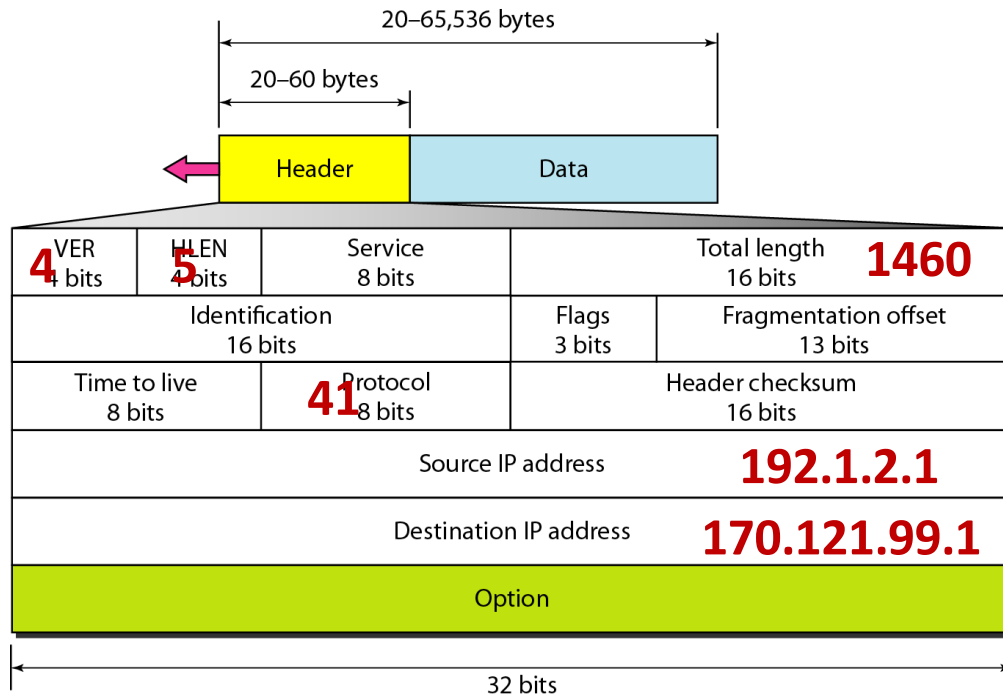


```
interface tunnel0
 tunnel mode ipv6ip
 tunnel source 192.1.2.1
 tunnel destination 170.121.99.1
 ipv6 address 2001:db80:e100:2::1/64
```

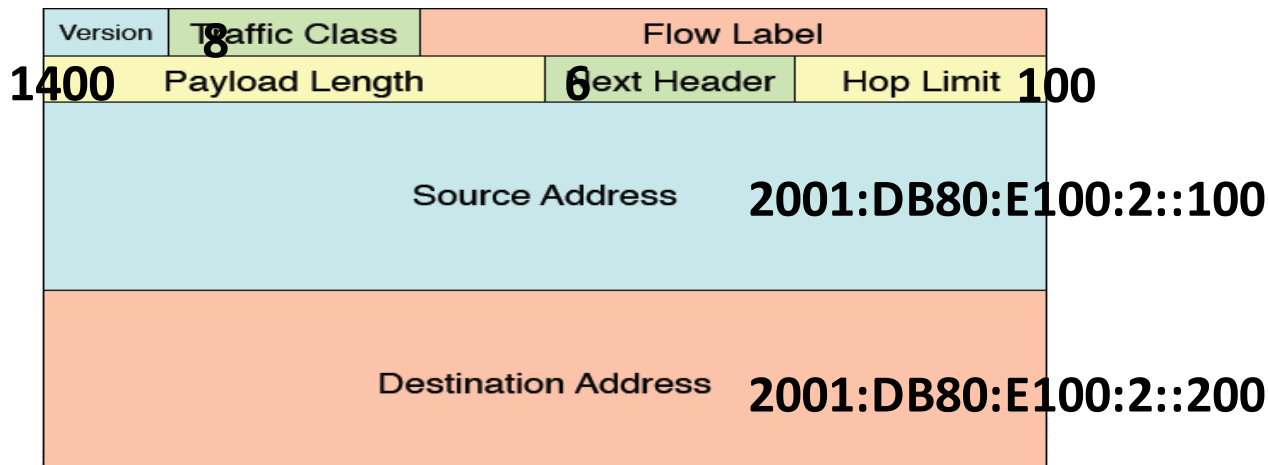
```
interface tunnel0
 tunnel mode ipv6ip
 tunnel source 170.121.99.1
 tunnel destination 192.1.2.1
 ipv6 address 2001:db80:e100:2::2/64
```

2001:db80:e100:2::/64 → tunnel0

Encapsulation



TTL from inner or configured
TOS from inner or configured



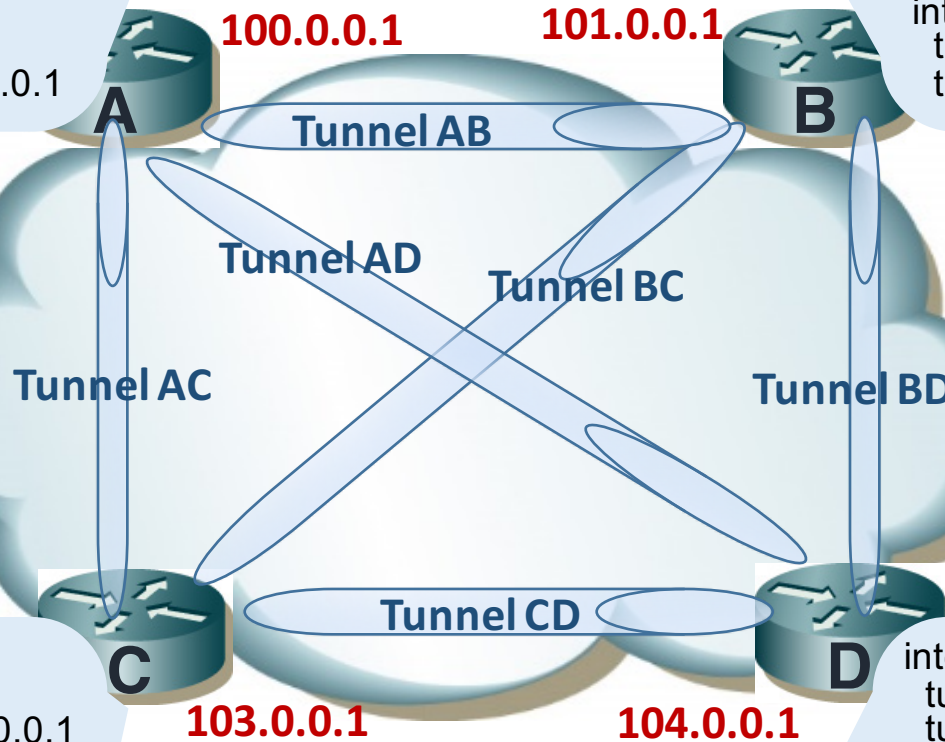
Potential Tunnel Issues

- MTU fragmentation
- ICMPv4 error handling
- Filtering protocol 41
- NAT (Network Address Translation)

```
interface tunnelAB
 tunnel source 100.0.0.1
 tunnel destination 101.0.0.1
!
interface tunnelAC
 tunnel source 100.0.0.1
 tunnel destination 103.0.0.1
!
interface tunnelAD
 tunnel source 100.0.0.1
 tunnel destination 104.0.0.1
```

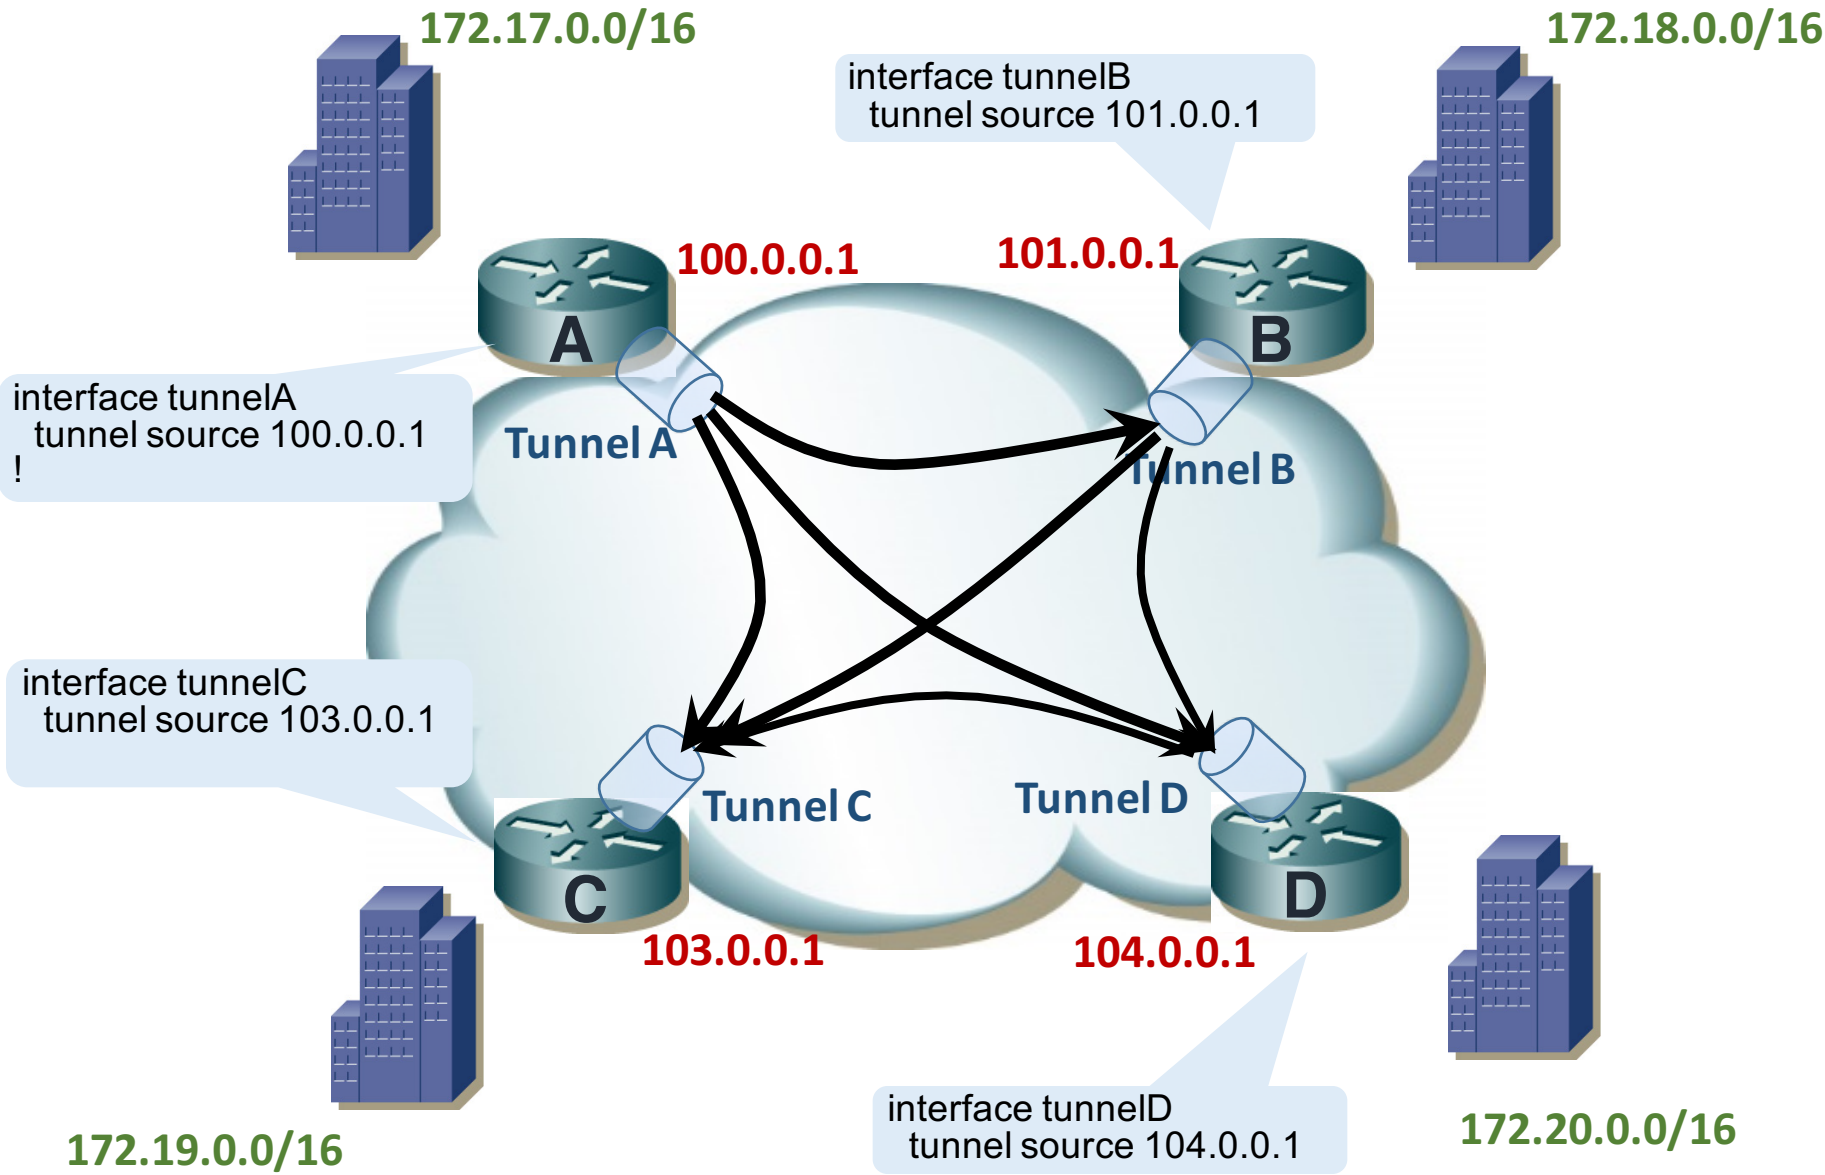
0.0/16

```
interface tunnelBA
 tunnel source 101.0.0.1
 tunnel destination 100.0.0.1
!
interface tunnelBC
 tunnel source 101.0.0.1
 tunnel destination 103.0.0.1
!
interface tunnelBD
 tunnel source 101.0.0.1
 tunnel destination 104.0.0.1
```



```
interface tunnelCA
 tunnel source 103.0.0.1
 tunnel destination 100.0.0.1
!
interface tunnelCB
 tunnel source 103.0.0.1
 tunnel destination 101.0.0.1
!
interface tunnelCD
 tunnel source 103.0.0.1
 tunnel destination 104.0.0.1
```

```
interface tunnelDA
 tunnel source 104.0.0.1
 tunnel destination 100.0.0.1
!
interface tunnelDC
 tunnel source 104.0.0.1
 tunnel destination 103.0.0.1
!
interface tunnelDB
 tunnel source 104.0.0.1
 tunnel destination 101.0.0.1
```



ISATAP

- **ISATAP** (Intra-Site Automatic Tunneling Addressing Protocol)
an automatic tunneling mechanism used inside an organization that has an IPv4-dominant backbone, but has selected users that need IPv6 capability

ISATAP Functions

- ISATAP connects dual-stack nodes, isolated within an IPv4-only network
 - To exchange IPv6 traffic with each other (host ISATAP)
 - To exchange traffic with the global IPv6 Internet
- ISATAP is a mechanism with minimal configuration required
- ISATAP is ideal when there are relatively few, relatively scattered individual nodes that need service

Link-Local ISATAP

192.0.2.100

IPv4 Address

Is converted to hex
form

C000:0264

0000:5EFE

And pre-pended with the
ISATAP 32-bit link-local suffix

::0000:5EFE:C000:0264

FE80::/10

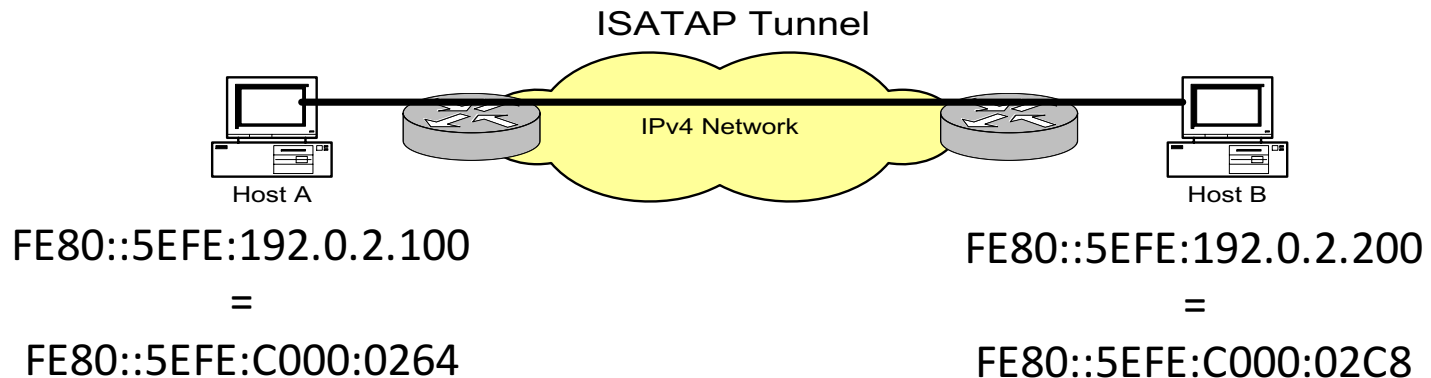
The link-local prefix merges with
the network identifier to create the

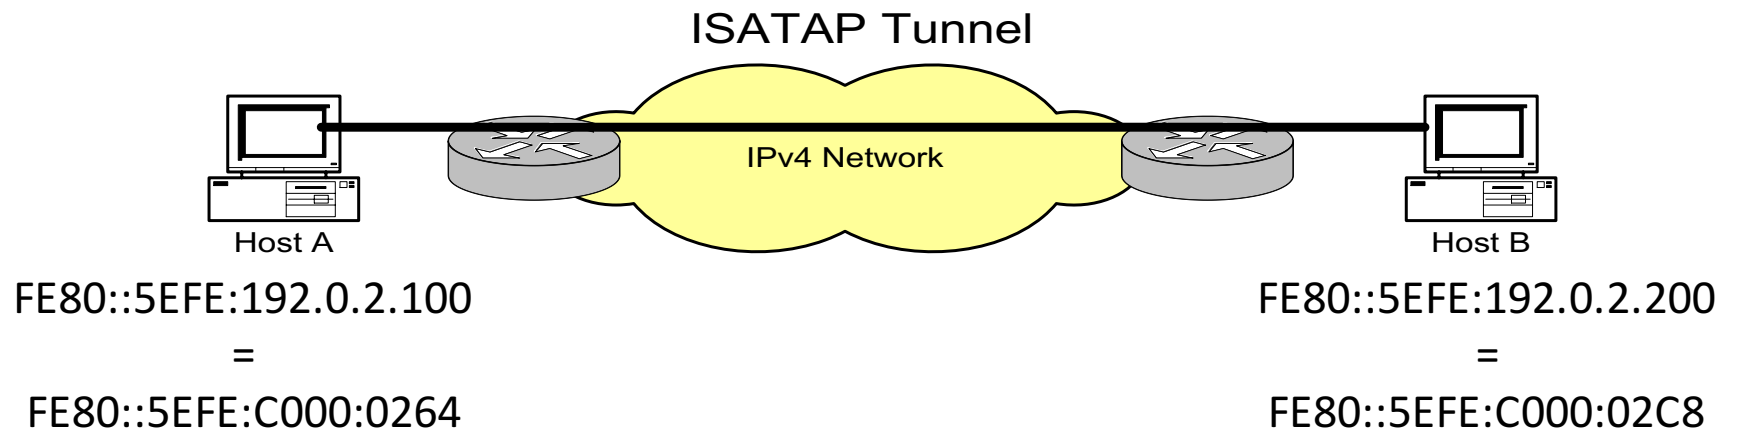
ISATAP IPv6 link-local address

FE80::0000:5EFE:C000:0264

Link-local ISATAP example

- Two ISATAP hosts exchanging packets using link-local addresses
- Only route on ISATAP hosts is “send all IPv6 traffic via ISATAP pseudo-IF”
- Hosts are many IPv4 hops away which appear link-local to IPv6





src: 192.0.1.100

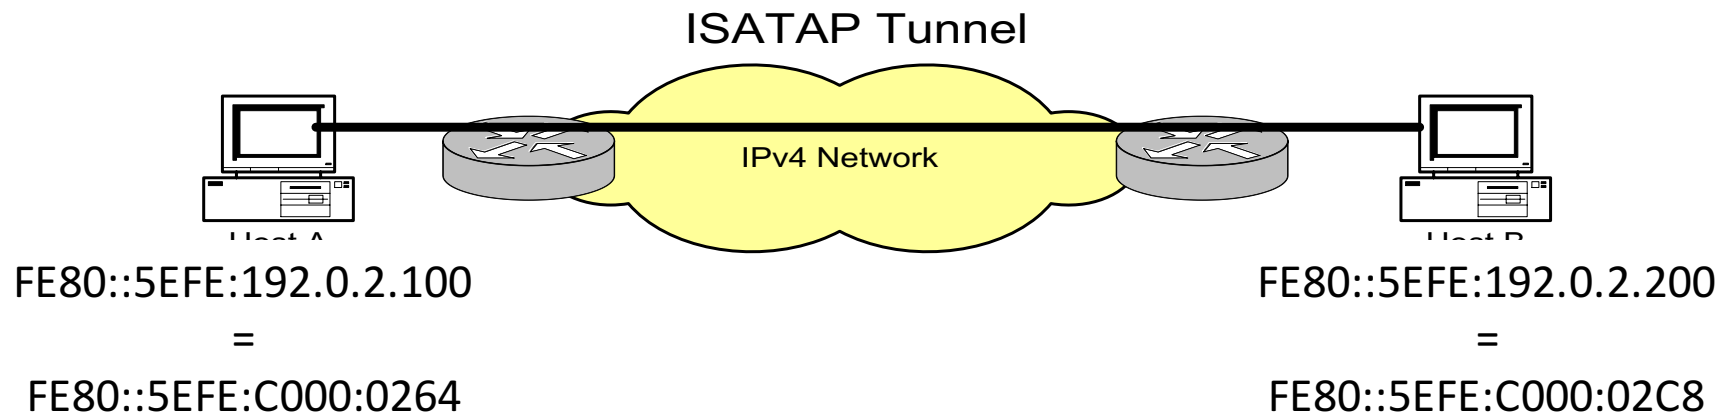
dst: 192.0.2.200

src: FE80::5EFE:C000:0264

dst: FE80::5EFE:C000:02C8



FE80::5EFE:192.0.2.10
=
FE80::5EFE:C000:020A

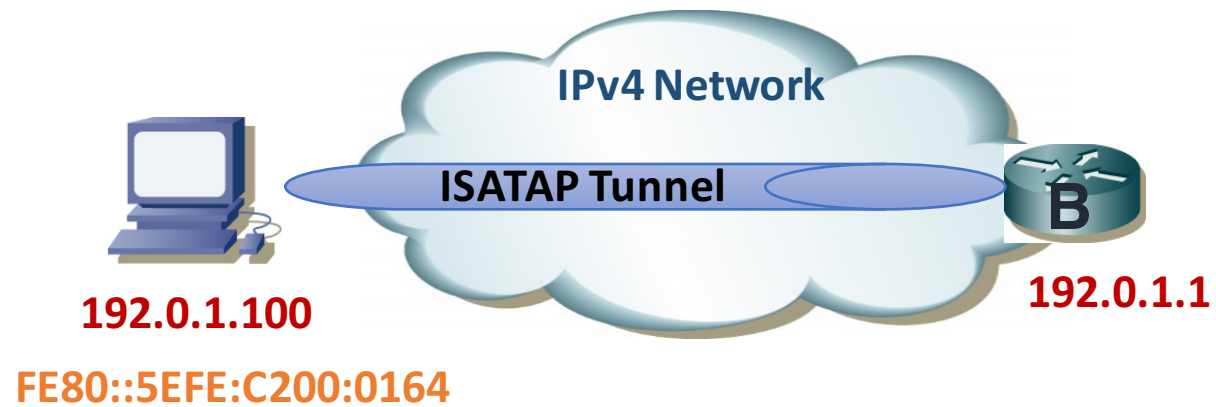


src: 192.0.1.100


dst: 192.0.2.10

src: FE80::5EFE:C000:0264

dst: FE80::5EFE:C000:020A

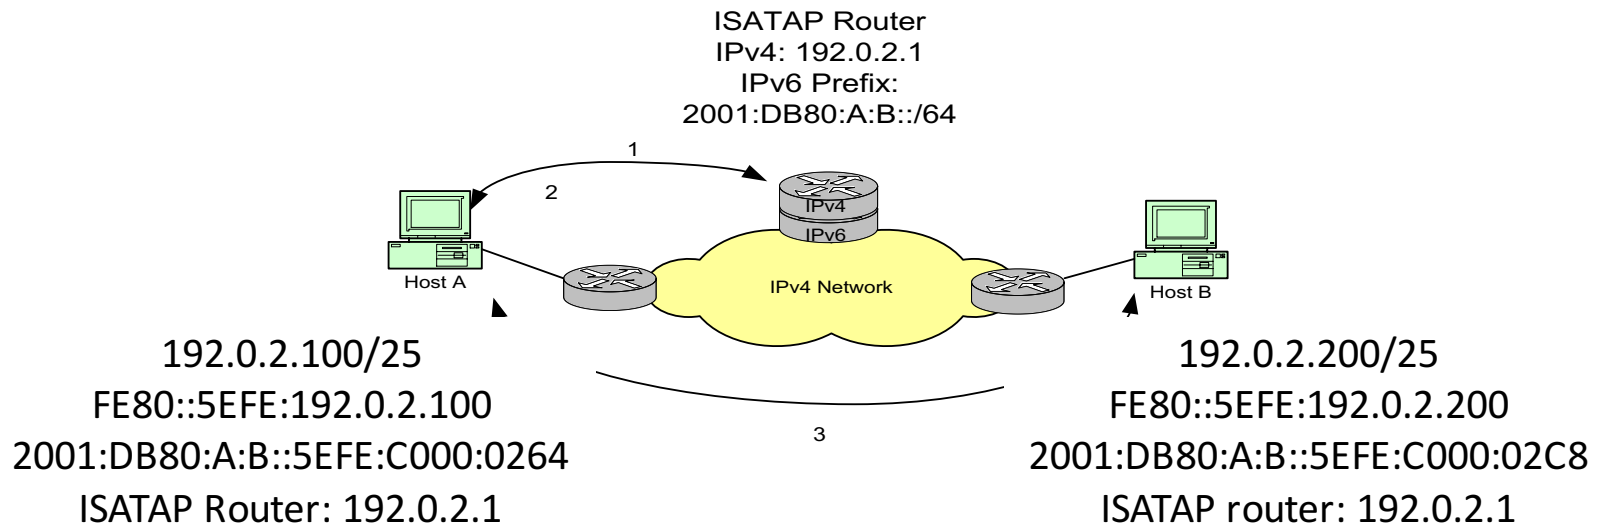


Src: FE80::5EFE:C200:0164
Dst: FE80::5EFE:C200:0101



Globally-routable ISATAP

- ISATAP more flexible when using an ISATAP router
- ISATAP hosts are configured with ISATAP router IPv4 address
- Hosts send router solicitation, inside tunnel, and ISATAP router responds



ISATAP Summary

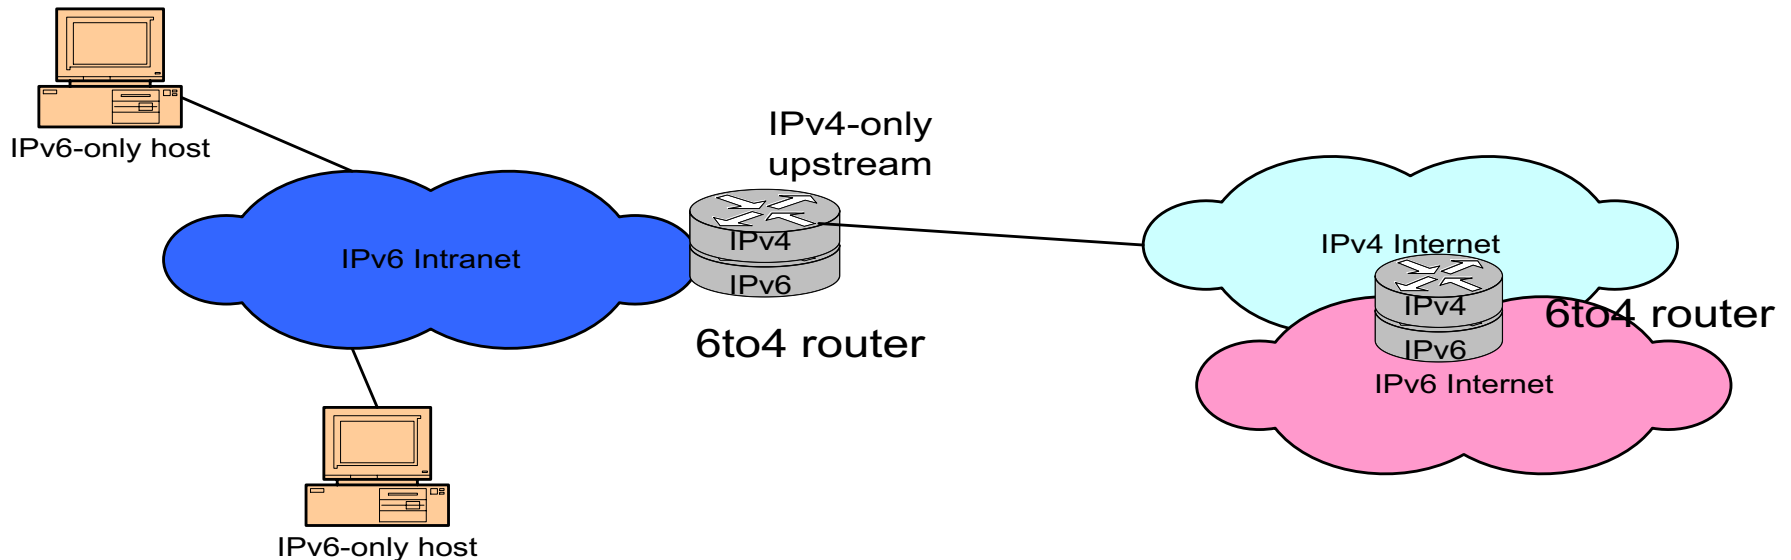
- ISATAP scales better than manually configured tunnels inside the enterprise
- Decapsulate-from-anywhere issues (like 6to4) mitigated by internal deployment
- No authentication provided – any dual stack node that knows ISATAP router address can obtain services
- May need to look at other alternatives if security is required

IPv6 6to4 Transition Mechanism

- 6to4 is an automatic tunneling mechanism that provides v6 capability to a dual-stack node or v6-capable site that has only IPv4 connectivity to the site

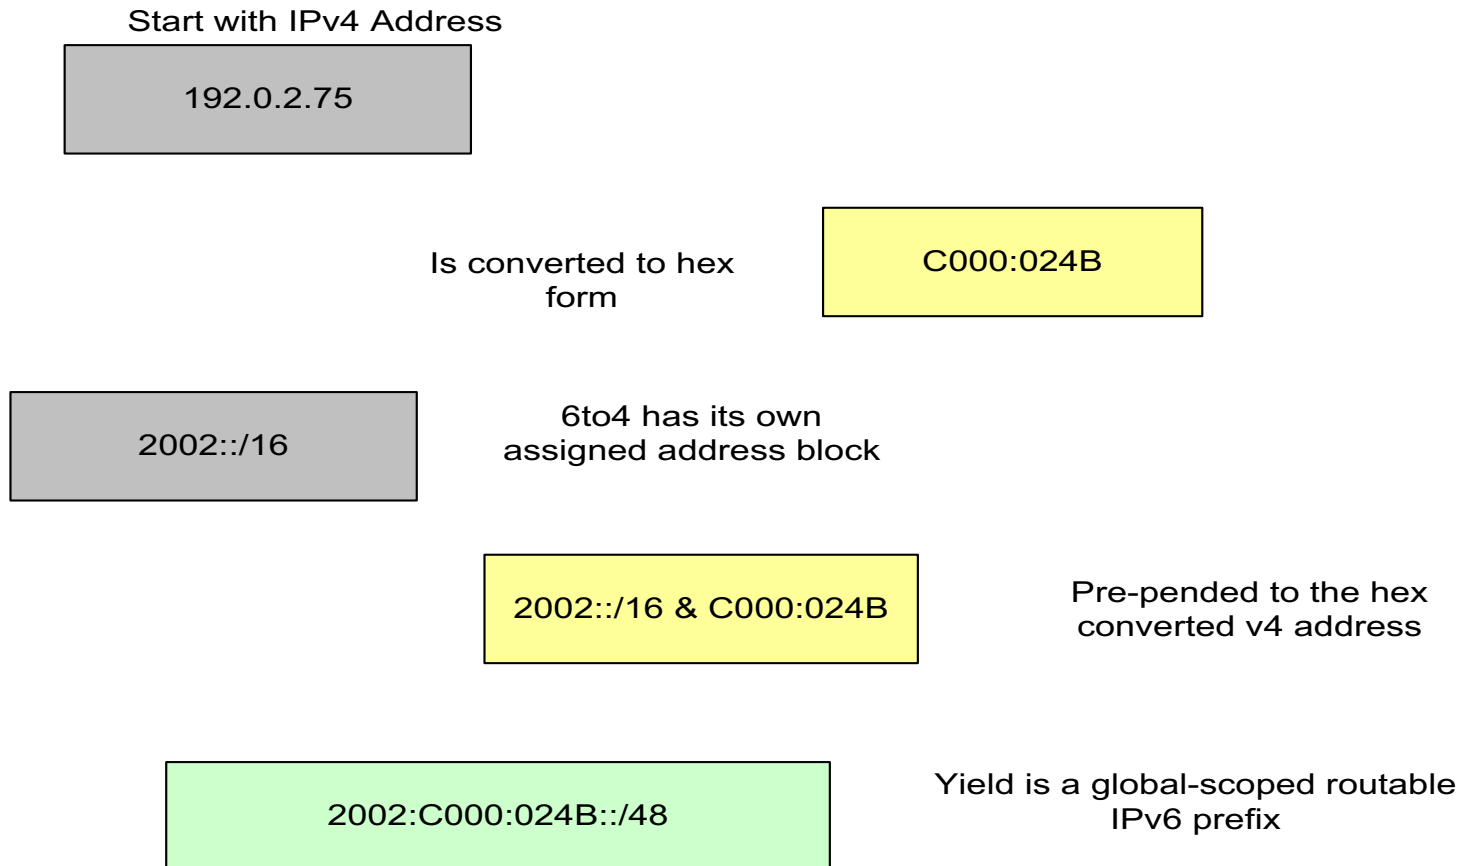
6to4 Basics

- 6to4 is an automatic tunnel mechanism
- Provides v6 upstream for v6-capable site over v4-only Internet connection
- Uses embedded addressing (v4addr embedded in v6addr) as do other automatic mechanisms

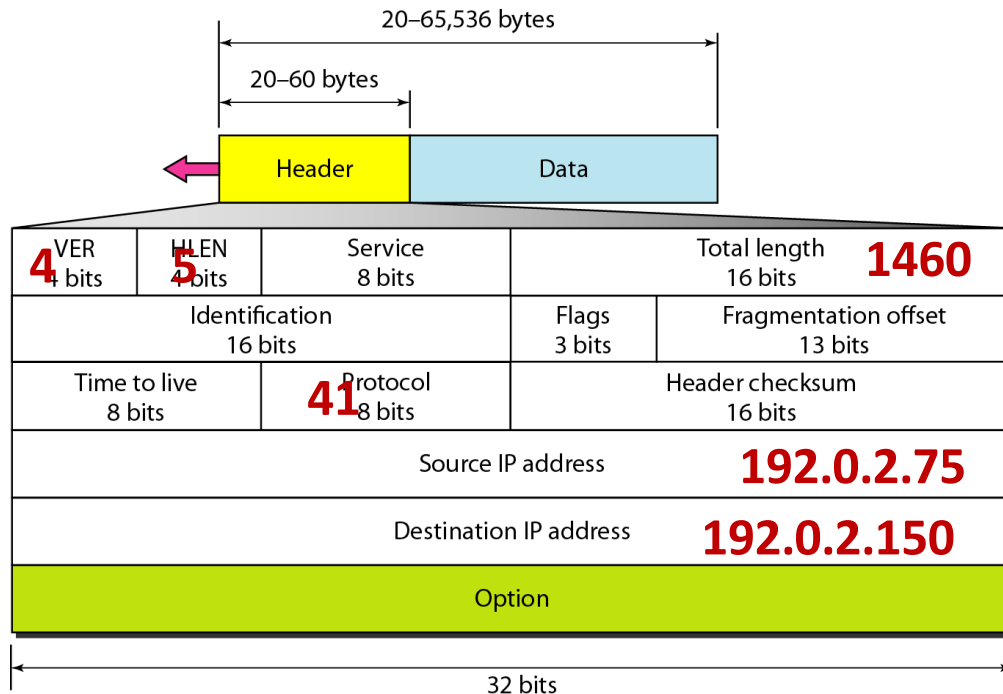


6to4 Address Construction

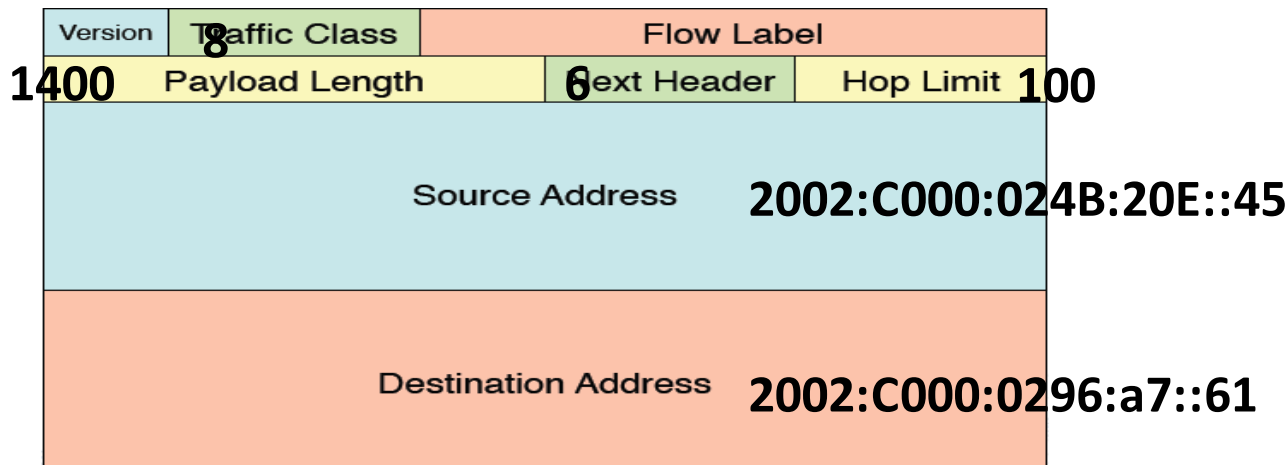
- 6to4 setups a valid, unique /48 IPv6 prefix from the outside IPv4 address of the site router

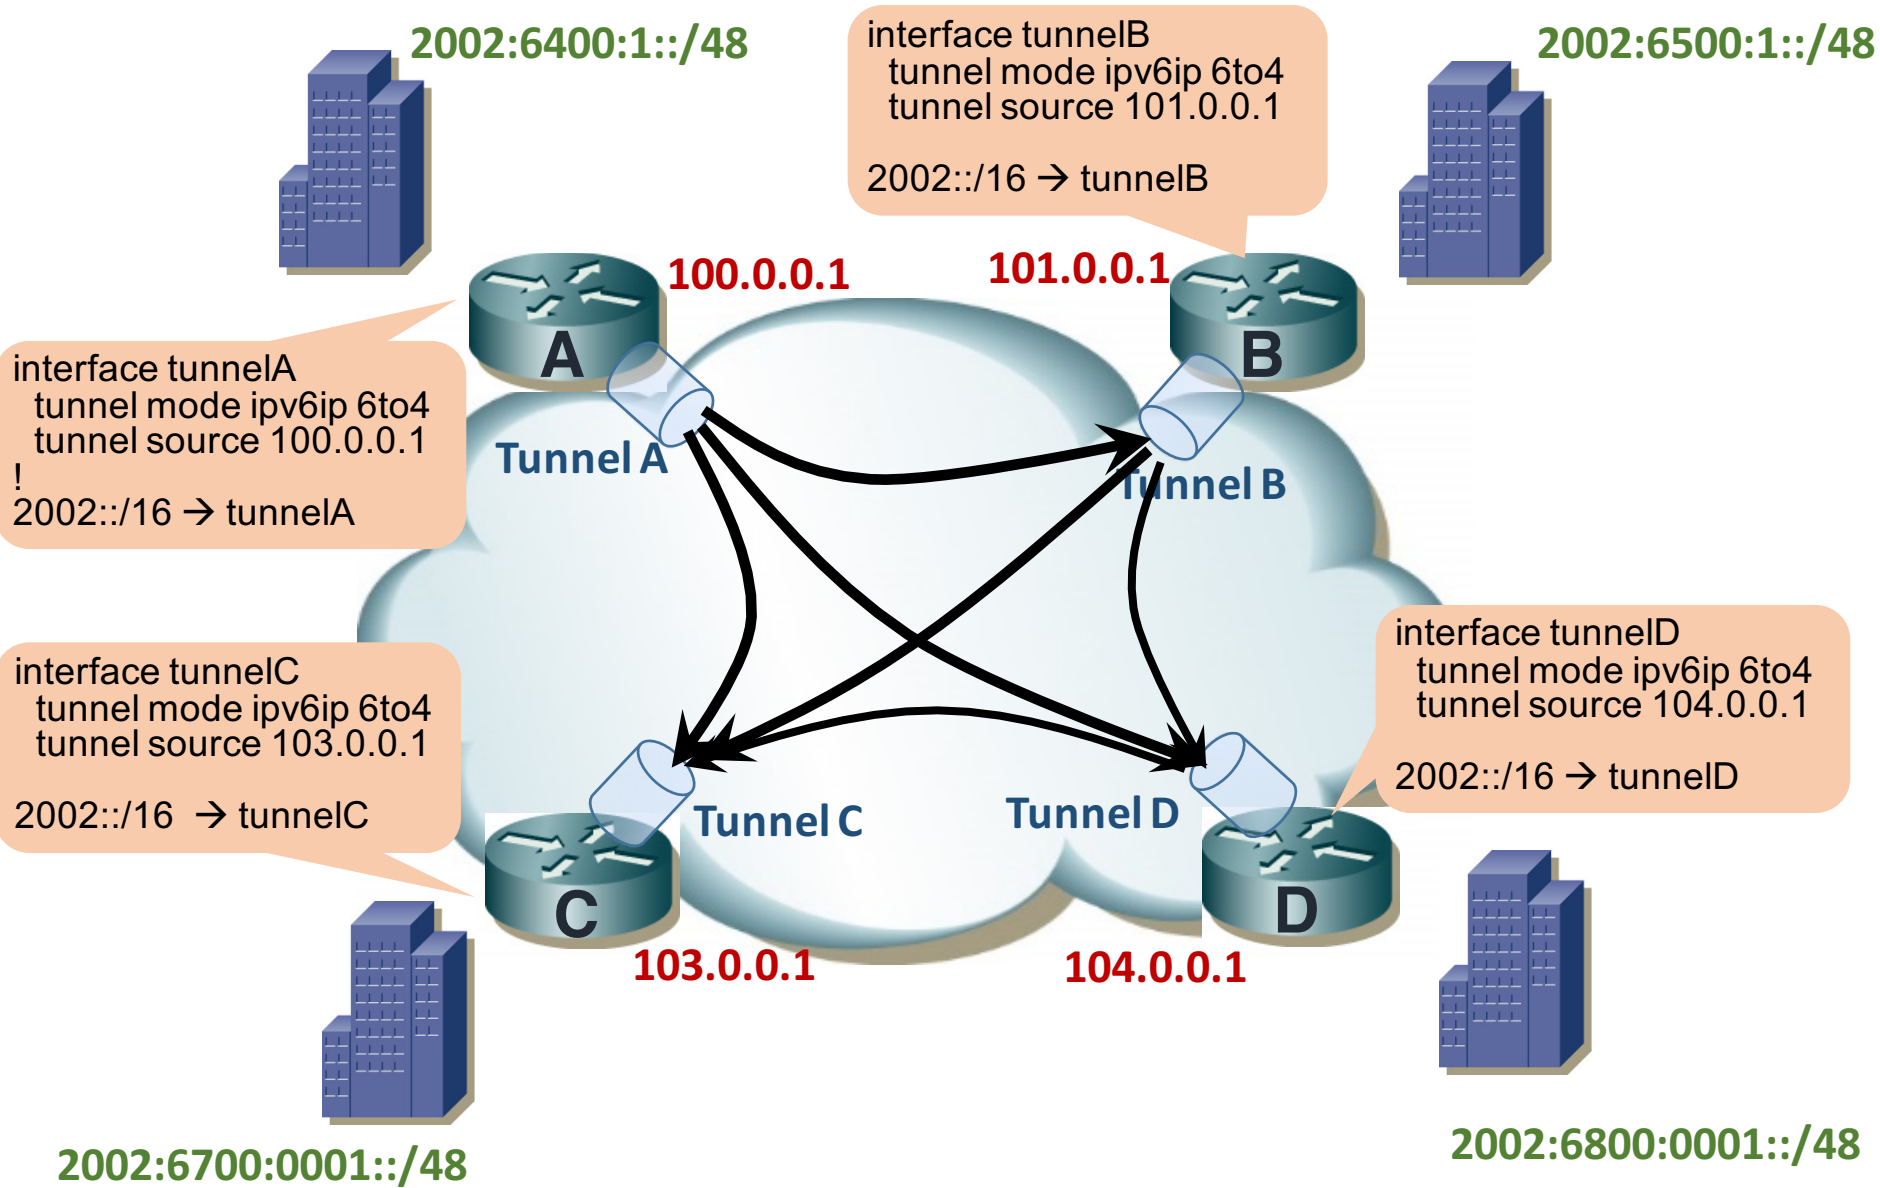


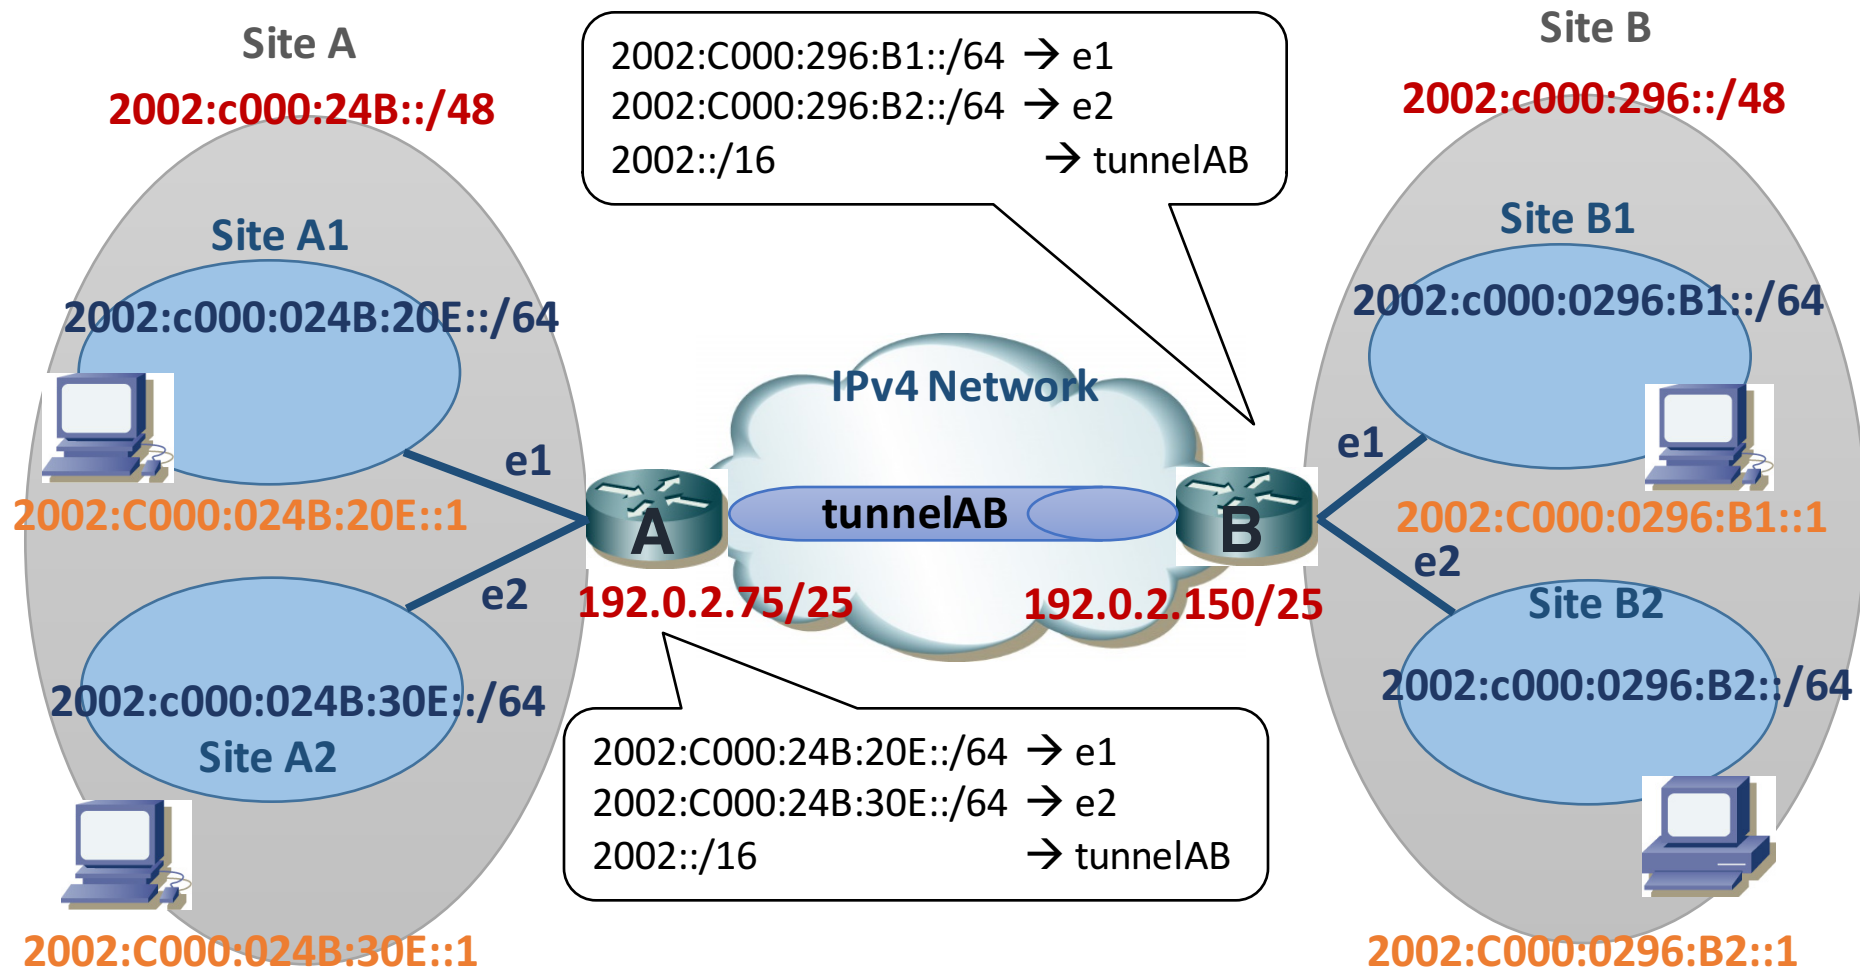
Encapsulation



TTL from inner or configured
TOS from inner or configured







Site A

Site A1



2002:C000:024B:20E::1

e1

Site A2



2002:C000:024B:30E::1

e2

2002:C000:296:B1::/64 → e1

2002:C000:296:B2::/64 → e2

2002::/16 → tunnelAB

IPv4 Network

tunnelAB

192.0.2.75/25

192.0.2.150/25

2002:C000:24B:20E::/64 → e1

2002:C000:24B:30E::/64 → e2

2002::/16 → tunnelAB

Site B

Site B1



2002:C000:0296:B1::1

e1

e2

Site B2



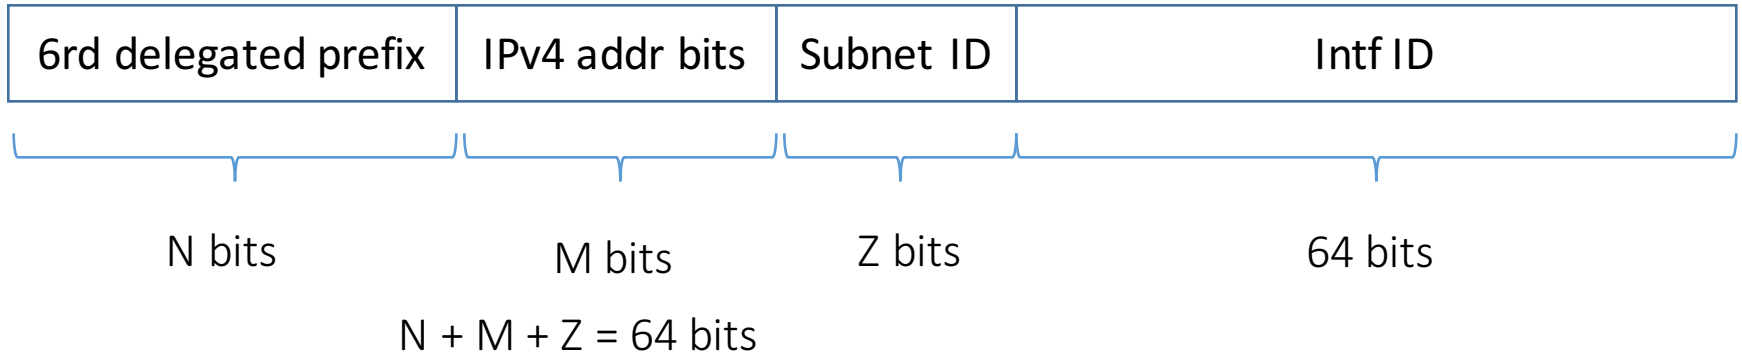
2002:C000:0296:B2::1

Src: 2002:C000:024B:30E::1
Src: 2002:C000:024B:30E::1
Dst: 2002:C000:024B:20E::1
Dst: 2002:C000:0296:B2::1

6RD Tunnel

- IPv6 Rapid Deployment on IPv4 Infrastructures (6rd) – RFC 5969
- Utilize an SP's own IPv6 address prefix rather than a well-known prefix (2002::/16)

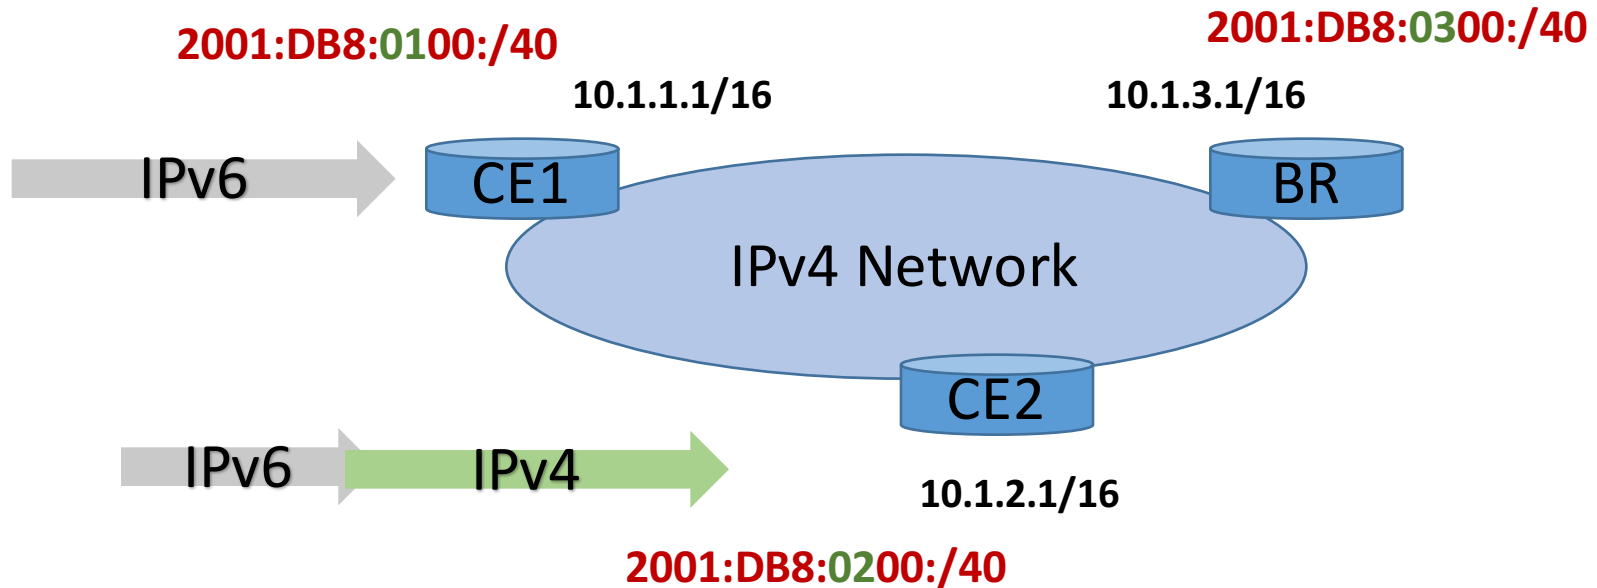
6RD Address Format



| | | |
|---------------|----------------|---------------|
| Common prefix | IPv4 addr bits | Common Suffix |
|---------------|----------------|---------------|

| Parameter | Value |
|---------------------------|---------------|
| 6rd Prefix/length | 2001:DB8::/32 |
| IPv4 Common prefix/length | 10.1.0.0/16 |
| IPv4 Common suffix/length | 0.0.0.1/8 |

6rd Prefix 2001:DB8::/32
IPv4 common prefix: 10.1.0.0/16
IPv4 common suffix: 0.0.0.1/8



IPv6: 2001:DB8:0100::C15C:0 → 2001:DB8:0200::C26B:0

IPv4: 10.1.1.1 → 10.1.2.1

6rd Prefix 2001:DB8::/32
IPv4 common prefix: 10.1.0.0/16
IPv4 common suffix: 0.0.0.1/8

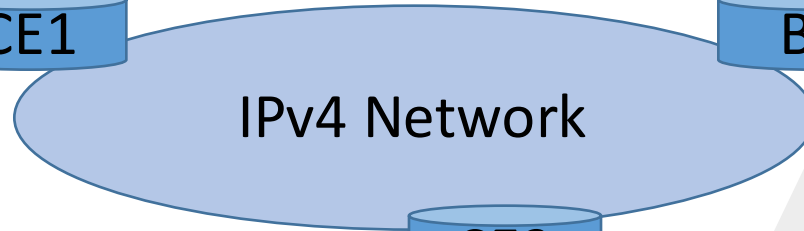
2001:DB8:0100:/40

10.1.1.1/16



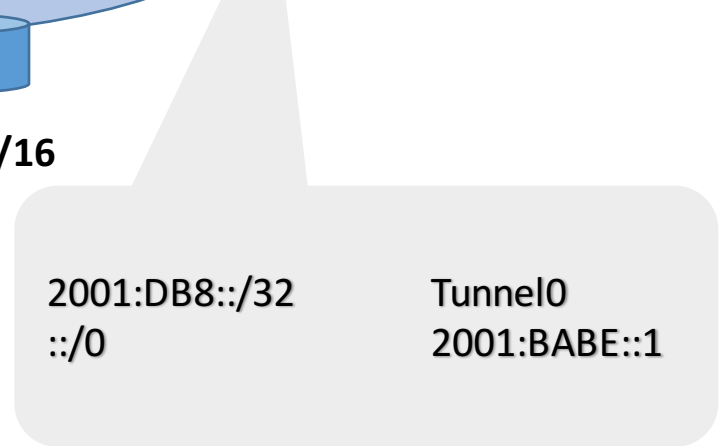
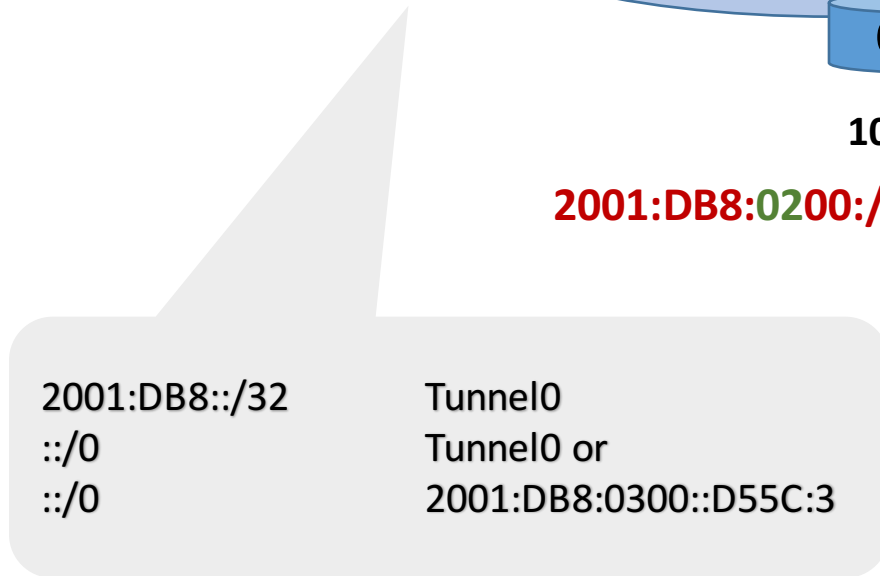
2001:DB8:0300:/40

10.1.3.1/16



10.1.2.1/16

2001:DB8:0200:/40



IPv6 Tunnel Address Format

